

Query Processing & Optimization

DB/AI Bootcamp

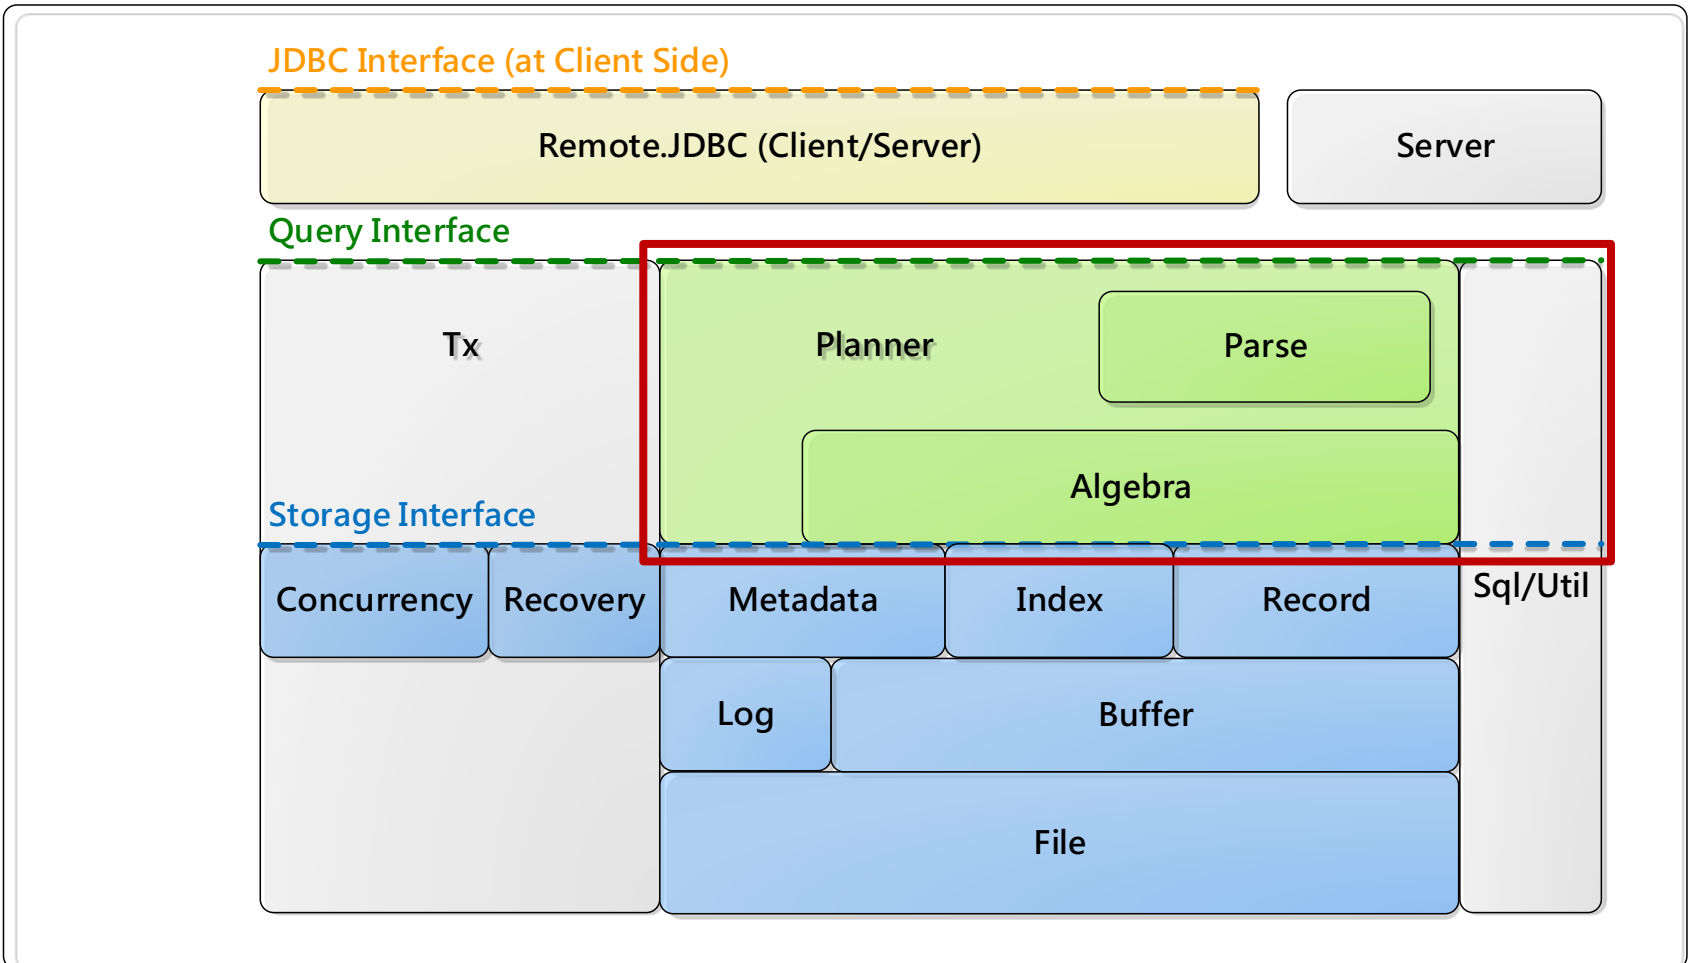
2018 Summer

DataLab, CS, NTHU

Query Processing

Query Engine

VanillaCore



Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - Predicates
 - Verifier
- Scans and plans
- Query planning
 - Deterministic planners

What does a DB do when SQL coming?

1. Parses the SQL command
2. Verifies the SQL command
3. Finds a good plan for the SQL command
4. Executes the plan

What does a DB do when SQL coming?

1. Parses the SQL command
2. Verifies the SQL command
3. Finds a good plan for the SQL command
4. Executes the plan

Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - Predicates
 - Verifier
- Scans and plans
- Query planning
 - Deterministic planners

SQL Statement Processing

- Input:
 - A SQL statement
- Output:
 - Internal *SQL data* object that can be fed to the constructors of various plans/scans
- Two stages:
 - *Parsing* (syntax-based)
 - *Verification* (semantic-based)

Syntax vs. Semantics

- The ***syntax*** of a language is a set of rules that describes the strings that could possibly be meaningful statements

- Is this statement syntactically legal?

```
SELECT FROM TABLES t1 AND t2 WHERE b - 3
```

- No
 - SELECT clause must refer to some field
 - TABLES is not a keyword
 - AND should separate predicates not tables
 - b - 3 is not a predicate

Syntax vs. Semantics

- Is this statement syntactically legal?
`SELECT a FROM t1, t2 WHERE b = 3`
 - Yes, we can infer that this statement is a query
 - But is it actually meaningful?
- The ***semantics*** of a language specifies the actual meaning of a syntactically correct string
- Whether it is semantically legal depends on
 - Is `a` a field name?
 - Are `t1`, `t2` the names of tables?
 - Is `b` the name of a numeric field?
- Semantic information is stored in the database's metadata (catalog)

Syntax vs. Semantics in VanillaCore

- `Parser` converts a SQL statement to SQL data based on the syntax
 - Exceptions are thrown upon syntax error
 - Outputs SQL data, e.g., `QueryData`, `InsertData`, `ModifyData`, `CreatTableData`, etc.
 - All defined in `query.parse` package
- `Verifier` examines the metadata to validate the semantics of SQL data
 - Defined in `query.planner` package

Outline

- Overview
- Scans and plans
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - **Lexer**, parser, and SQL data
 - Predicates
 - Verifier
- Query planning
 - Deterministic planners

Parsing SQL Commands

- Parser uses a *parsing algorithm* to convert a SQL string to SQL data
 - To be detailed later
- Uses a *lexical analyzer* (also called *lexer* or tokenizer) that splits the SQL string into tokens when reading

SELECT a FROM t1, t2 WHERE b = 3

Stream-based API

- Reads a SQL string only *once*
- `matchXXX`
 - Returns whether the next token is of the specified type
- `eatXXX`
 - Returns the value of the next token if the token is of the specified type
 - Otherwise throws `BadSyntaxException`

Lexer
- keywords : Collection<String> - tok : StreamTokenizer
+ Lexer(s : String) + matchDelim(delimiter : char) : boolean + matchNumericConstant() : boolean + matchStringConstant() : boolean + matchKeyword(keyword : String) : boolean + matchId() : boolean + eatDelim(delimiter : char) + eatNumericConstant() : double + eatStringConstant() : String + eatKeyword(keyword : String) + eatId() : String

Outline

- Overview
- Scans and plans
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, **parser**, and SQL data
 - Predicates
 - Verifier
- Query planning
 - Deterministic planners

Grammar

- A ***grammar*** is a set of rules that describe how tokens can be legally combined
- E.g.,

<Field>	:=	IdTok
<Constant>	:=	StrTok NumericTok
<Expression>	:=	<Field> <Constant>
<Term>	:=	<Expression> = <Expression>
<Predicate>	:=	<Term> [AND <Predicate>]

SQL Data

- Parser returns SQL data
 - E.g., when the parsing the query statement (syntactic category `<Query>`), parser will returns a `QueryData` object
- All SQL data are defined in `query.parse` package

Parser and QueryData

Parser
- lex : Lexer
+ Parser(s : String) + updateCmd() : Object + query() : QueryData - id() : String - constant() : Constant - queryExpression() : Expression - term() : Term - predicate() : Predicate ... - create() : Object - delete() : DeleteData - insert() : InsertData - modify() : ModifyData - createTable() : CreateTableData - createView() : CreateViewData - createIndex() : CreateIndexData

QueryData
+ QueryData(projFields : Set<String>, tables : Set<String>, pred : Predicate, groupFields : Set<String>, aggFn : Set<AggregationFn>, sortFields : List<String>, sortDirs : List<Integer>) + projectFields() : Set<String> + tables() : Set<String> + pred() : Predicate + groupFields() : Set<String> + aggregationFn() : Set<String> + sortFields() : List<String> + sortDirs() : List<Integer> + toString() : String

Other SQL data

InsertData
<ul style="list-style-type: none">+ InsertData(tblname : String, flds : List<String>, vals : List<Constant>)+ tableName() : String+ fields() : List<String>+ val() : List<Constant>

CreateTableData
<ul style="list-style-type: none">+ InsertData(tblname : String, sch : Schema)+ tableName() : String+ newSchema : Schema

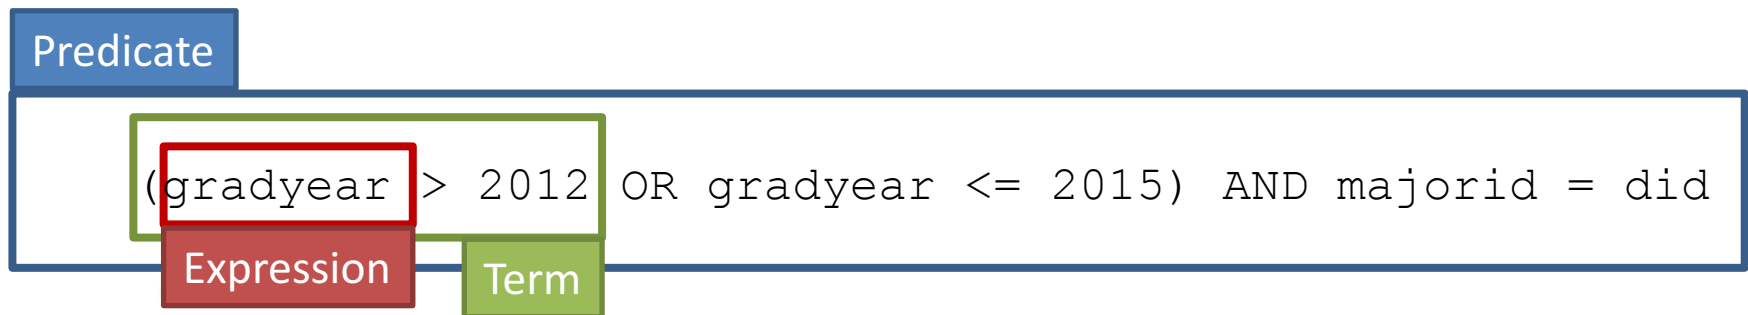
Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - **Predicates**
 - Verifier
- Scans and plans
- Query planning
 - Deterministic planners

Predicate

```
<Field>      := IdTok  
<Constant>   := StrTok | NumericTok  
<Expression> := <Field> | <Constant>  
<Term>        := <Expression> = <Expression>  
<Predicate>  := <Term> [ AND <Predicate> ]
```

- Classes defined in `sql.predicates` in `VanillaCore`
- For example,




Term

- **Term supports five operators**
 - `OP_EQ (=)`, `OP_LT (<)`, `OP_LTE (<=)`,
`OP_GE (>)`, and `OP_GTE (>=)`

Term
<u><<final>> + OP_EQ : Operator</u> <u><<final>> + OP_LT : Operator</u> <u><<final>> + OP_LTE : Operator</u> <u><<final>> + OP_GE : Operator</u> <u><<final>> + OP_GTE : Operator</u>
+ Term(lhs : Expression, op : Operator, rhs : Expression) + operator(fldname : String) : Operator + oppositeConstant(fldname : String) : Constant + oppositeField(fldname : String) : String + isApplicableTo(sch : Schema) : boolean + isSatisfied(rec : Record) : boolean + toString() : String

<u><<abstract>></u> <u>Operator</u>
<<abstract>> complement() : Operator <<abstract>> isSatisfied(lhs : Expression, rhs : Expression, rec : Record) : boolean



Methods of Term

- The method `isSatisfied(rec)` returns true if given the specified record, the two expressions evaluate to matching values

Term5: `created = 2012/11/15`

	blog_id	url	created	author_id
X	33981	...	2009/10/31	729
O	33982	...	2012/11/15	730
X	41770	...	2012/10/20	729

```
public boolean isSatisfied(Record rec) {  
    return op.isSatisfied(lhs, rhs, rec);  
}
```

Predicate

- A predicate in VanillaCore is a conjunct of terms, e.g., *term1 AND term2 AND ...*

Predicate
<pre>+ Predicate() + Predicate(t : Term) // used by the parser + conjunctWith(t : Term) // used by a scan + isSatisfied(rec : Record) : boolean // used by the query planner + selectPredicate(sch : Schema) : Predicate + joinPredicate(sch1 : Schema, sch2 : Schema) : Predicate + constantRange(fldname : String) : ConstantRange + joinFields(fldname : String) : Set<String> + toString() : String</pre>

Creating a Predicate in a Query Parser

```
// majorid <=30 AND majorid=did
Expression exp1 = new FieldNameExpression("majorid");
Expression exp2 = new ConstantExpression(
    new IntegerConstant(30));
Term t1 = new Term(exp1, OP_LTE, exp2);

Expression exp3 = new FieldNameExpression("majorid");
Expression exp4 = new FieldNameExpression("did");
Term t2 = new Term(exp3, OP_EQ, exp4);

Predicate pred = new Predicate(t1);
pred.conjunctWith(t2);
```

Outline

- Overview
- Scans and plans
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - Predicates
 - **Verifier**
- Query planning
 - Deterministic planners

Verification

- Before feeding the SQL data into the plans/scans, the planner asks the `Verifier` to verify the semantics correctness of the data

Verification

- The `Verifier` checks whether:
 - The mentioned tables and fields actually exist in the catalog
 - The mentioned fields are not ambiguous
 - The actions on fields are type-correct
 - All constants are of correct type and size to their corresponding fields

Verifying the INSERT Statement

```
public static void verifyInsertData(InsertData data, Transaction tx) {  
    // examine table name  
    TableInfo ti = VanillaDb.catalogMgr().getTableInfo(data.tableName(), tx);  
    if (ti == null)  
        throw new BadSemanticException("table " + data.tableName() + " does not exist");  
  
    Schema sch = ti.schema();  
    List<String> fields = data.fields();  
    List<Constant> vals = data.vals();  
  
    // examine whether values have the same size with fields  
    if (fields.size() != vals.size())  
        throw new BadSemanticException("#fields and #values does not match");  
  
    // verify field existence and type  
    for (int i = 0; i < fields.size(); i++) {  
        String field = fields.get(i);  
        Constant val = vals.get(i);  
        // check field existence  
        if (!sch.hasField(field))  
            throw new BadSemanticException("field " + field + " does not exist");  
        // check whether field matches value type  
        if (!verifyConstantType(sch, field, val))  
            throw new BadSemanticException("field " + field  
                + " doesn't match corresponding value in type");  
    }  
}
```

Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Predicates
 - Lexer, parser, and SQL data
 - Verifier
- **Scans and plans**
- Query planning
 - Deterministic planners

What does a DB do when SQL coming?

1. Parses the SQL command
2. Verifies the SQL command
3. Finds a good *plan* for the SQL command
4. Executes the plan


Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Predicates
 - Lexer, parser, and SQL data
 - Verifier
- **Scans and plans**
- Query planning
 - Deterministic planners

SQL and Relational Algebra (1/2)

- A SQL command will be expressed as at-least one tree in relational algebra


```
SELECT b.blog_id
FROM blog_pages b, users u
WHERE b.author_id=u.user_id
      AND u.name='Steven Sinofsky'
      AND b.created >= 2011/1/1;
```



Why this translation?


SQL and Relational Algebra (2/2)

- SQL is difficult to implement directly
 - A single SQL command can embody several tasks
- Relational algebra is relatively easy to implement
 - Each **operator** denotes a small, well-defined task



Operators


- Single-table operators
 - select, project, sort, rename, extend, groupby, etc.
- Two-table operators
 - product, join, semijoin, etc.
- Operands
 - Tables, views, output of other operators, predicates, etc.
- Output
 - Always a table
 - To be returned or used as a param of the next op



Algebra Tree

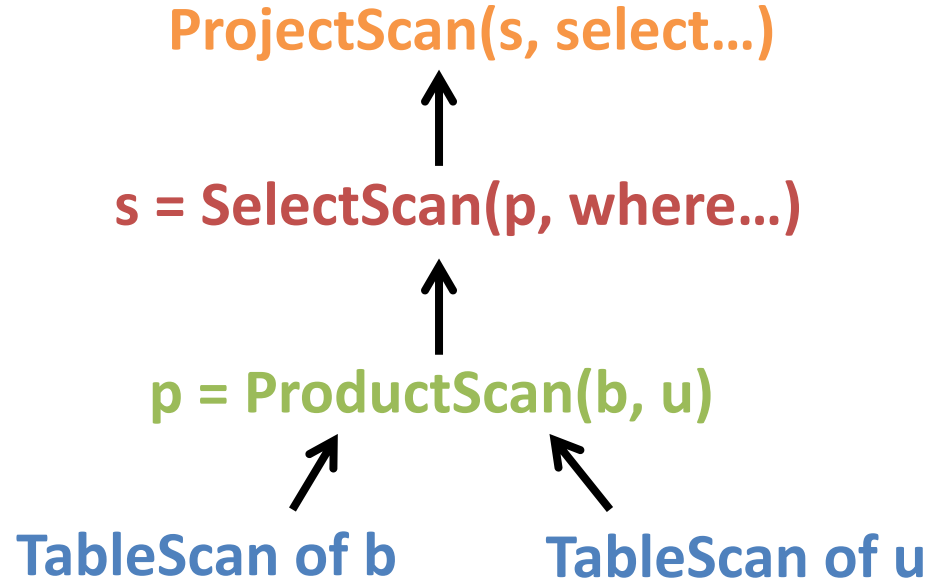
- A SQL command will be expressed as at-least one tree in relational algebra

```
SELECT b.blog_id
FROM blog_pages b, users u
WHERE b.author_id=u.user_id
      AND u.name='Steven Sinofsky'
      AND b.created >= 2011/1/1;
```




Scans

- A **scan** represents an operator in a relational algebra tree



The Scan Interface

- An iterator of output records of a partial query



Using a Scan

```
public static void printNameAndGradyear(Scan s) {  
    s.beforeFirst();  
    while (s.next()) {  
        Constant sname = s.getVal("sname");  
        Constant gradyear = s.getVal("gradyear");  
        System.out.println(sname + "\t" + gradyear);  
    }  
    s.close();  
}
```


Basic Scans

```
public SelectScan(Scan s, Predicate pred);
```

```
public ProjectScan(Scan s,  
    Collection<String> fieldList);
```

```
public ProductScan(Scan s1, Scan s2);
```

```
public TableScan(TableInfo ti, Transaction tx);
```


Building a Scan Tree

```
VanillaDb.init("studentdb");
Transaction tx =
    VanillaDb.txMgr().transaction(
        Connection.TRANSACTION_SERIALIZABLE, true);
TableInfo ti =
    VanillaDb.catalogMgr().getTableInfo("b", tx);

Scan ts = new TableScan(ti, tx);
Predicate pred = new Predicate("..."); // sid = 5

Scan ss = new SelectScan(ts, pred);
Collection<String> projectFld =
    Arrays.asList("sname");
Scan ps = new ProjectScan(ss, projectFld);

ps.beforeFirst();
while (ps.next())
    System.out.println(ps.getVal("sname"));
ps.close();
```



```

public class TableScan implements UpdateScan {
    private RecordFile rf;
    private Schema schema;

    public TableScan(TableInfo ti, Transaction tx) {
        rf = ti.open(tx);
        schema = ti.schema();
    }

    public void beforeFirst() {
        rf.beforeFirst();
    }

    public boolean next() {
        return rf.next();
    }

    public void close() {
        rf.close();
    }

    public Constant getVal(String fldName) {
        return rf.getVal(fldName);
    }

    public boolean hasField(String fldName) {
        return schema.hasField(fldName);
    }

    public void setVal(String fldName, Constant val) {
        rf.setVal(fldName, val);
    }
    ...
}

```

TableScan

- Basically, tasks are delegated to a RecordFile

SelectScan

```
public class SelectScan implements UpdateScan {
    private Scan s;
    private Predicate pred;

    public SelectScan(Scan s, Predicate pred) {
        this.s = s;
        this.pred = pred;
    }

    public boolean next() {
        while (s.next())
            // if current record satisfied the predicate
            if (pred.isSatisfied(s))
                return true;
        return false;
    }

    public void setVal(String fldname, Constant val) {
        UpdateScan us = (UpdateScan) s;
        us.setVal(fldname, val);
    }
    ...
}
```

ProjectScan

```
public class ProjectScan implements Scan {
    private Scan s;
    private Collection<String> fieldList;

    public ProjectScan(Scan s, Collection<String> fieldList) {
        this.s = s;
        this.fieldList = fieldList;
    }

    public boolean next() {
        return s.next();
    }

    public Constant getVal(String fldName) {
        if (hasField(fldName))
            return s.getVal(fldName);
        else
            throw new RuntimeException("field " + fldName + " not found.");
    }
    ...
}
```

ProductScan

```
public class ProductScan implements Scan {
    private Scan s1, s2;
    private boolean isLhsEmpty;

    public ProductScan(Scan s1, Scan s2) {
        this.s1 = s1;
        this.s2 = s2;
        s1.beforeFirst();
        isLhsEmpty = !s1.next();
    }

    public boolean next() {
        if (isLhsEmpty)
            return false;
        if (s2.next())
            return true;
        else if (!(isLhsEmpty = !s1.next())) {
            s2.beforeFirst();
            return s2.next();
        } else
            return false;
    }

    public Constant getVal(String fldName) {
        if (s1.hasField(fldName))
            return s1.getVal(fldName);
        else
            return s2.getVal(fldName);
    }
    ...
}
```

- Iterates through records following the *nested loops*

Example

project(s, select blog_id)

↓ beforeFirst()

**select(p, where name = 'Picachu'
and author_id = user_id)**


↓ beforeFirst()

product(b, u)

beforeFirst()
next()


beforeFirst()

b



blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u




user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

```
SELECT blog_id FROM b, u  
WHERE name = "Picachu"  
AND author_id = user_id;
```

Example

project(s, select blog_id)



**select(p, where name = 'Picachu'
and author_id = user_id)**



product(b, u)

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	729	729	Steven Sinofsky	10,235

next()

b


blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Example

project(s, select blog_id)



select(p, where name = 'Picachu'
and author_id = user_id)



product(b, u)

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	729	730	Picachu	NULL

next()

b


blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Example


project(s, select blog_id)



**select(p, where name = 'Picachu'
and author_id = user_id)**



product(b, u)



b


blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Example

project(s, select blog_id)



**select(p, where name = 'Picachu'
and author_id = user_id)**



product(b, u)

blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	730	729	Steven Sinofsky	10,235

next()

b

blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729


u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Example

blog_id
33982

project(s, select blog_id)



blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	730	730	Picachu	NULL

**select(p, where name = 'Picachu'
and author_id = user_id)**



product(b, u)

blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	730	730	Picachu	NULL

next()

b

blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Example

project(s, select...)



**select(p, where
name = 'Picachu')**



product(b, u)

getVal()

blog_id
33982

blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	730	730	Picachu	NULL

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	729	729	Steven Sinofsky	10,235
33981	...	2009/10/31	729	730	Picachu	NULL
33982	...	2012/11/15	730	729	Steven Sinofsky	10,235
33982	...	2012/11/15	730	730	Picachu	NULL
41770	...	2012/10/20	729	729	Steven Sinofsky	10,235
41770	...	2012/10/20	729	730	Picachu	NULL

b

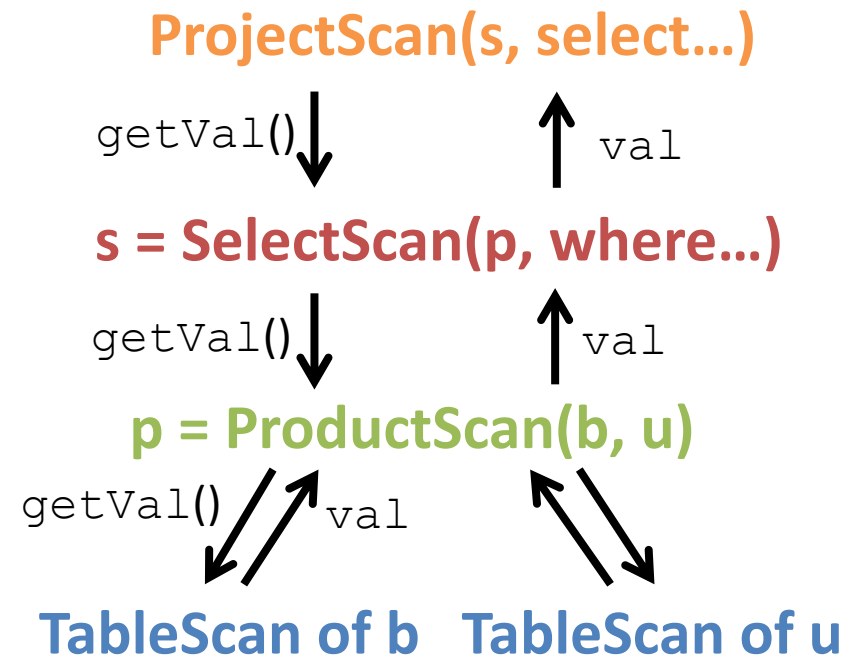
blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	730
41770	...	2012/10/20	729

u

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

Pipelined Scanning

- The above operators implement **pipelined scanning**
 - Calling a method of a node results in recursively calling the same methods of child nodes on-the-fly
 - Records are computed one at a time as needed---no intermediate records are saved



Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - Predicates
 - Verifier
- Scans and **plans**
- Query planning
 - Deterministic planners

Scan Tree for SQL Command?

- Given the scans:

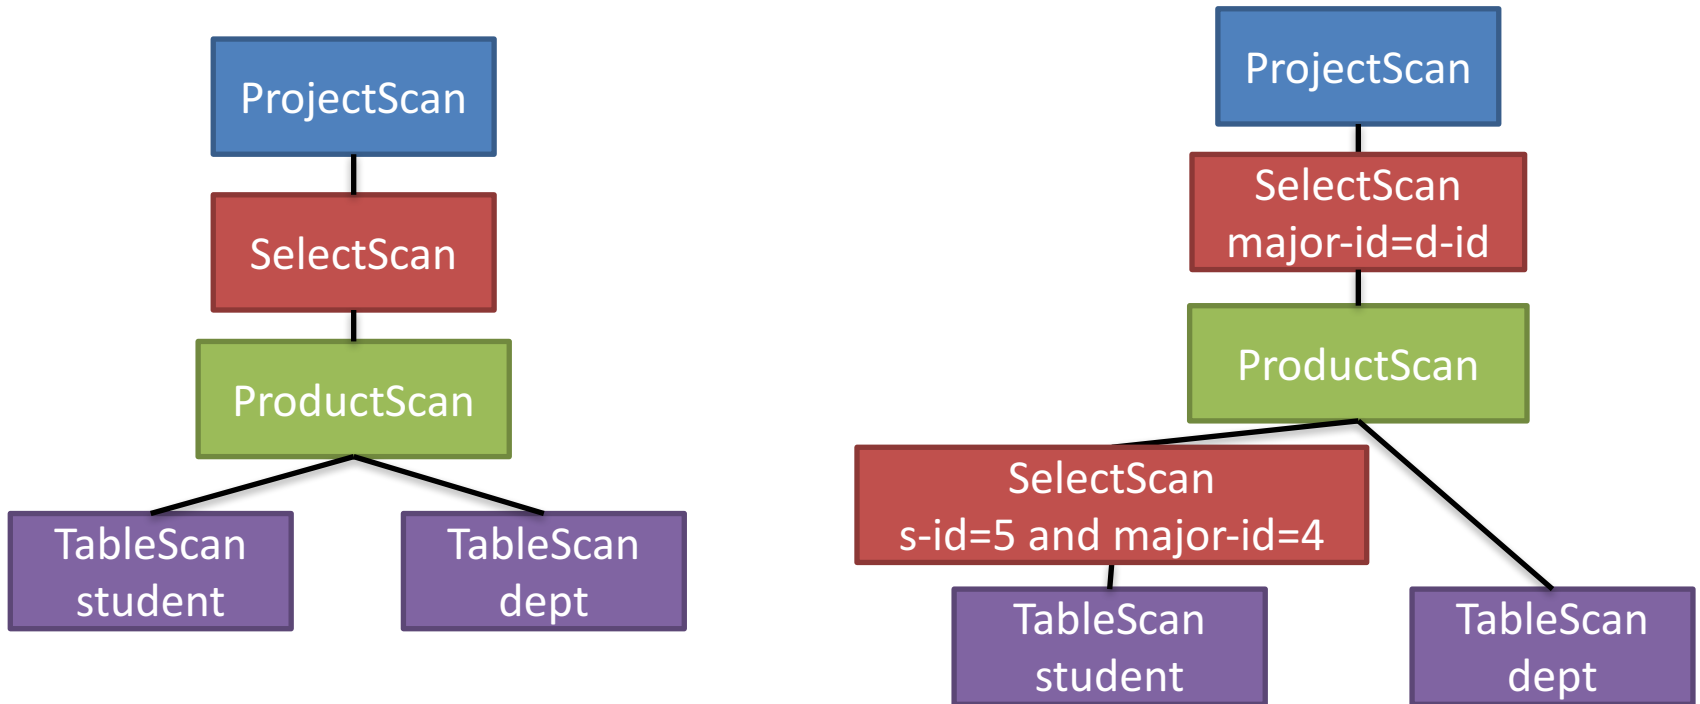


- Can you build a scan tree for this query:

```
SELECT sname FROM student, dept
      WHERE major-id = d-id
            AND s-id = 5 AND major-id = 4;
```


Which One is Better?

```
SELECT sname FROM student, dept
      WHERE major-id = d-id
            AND s-id = 5 AND major-id = 4;
```

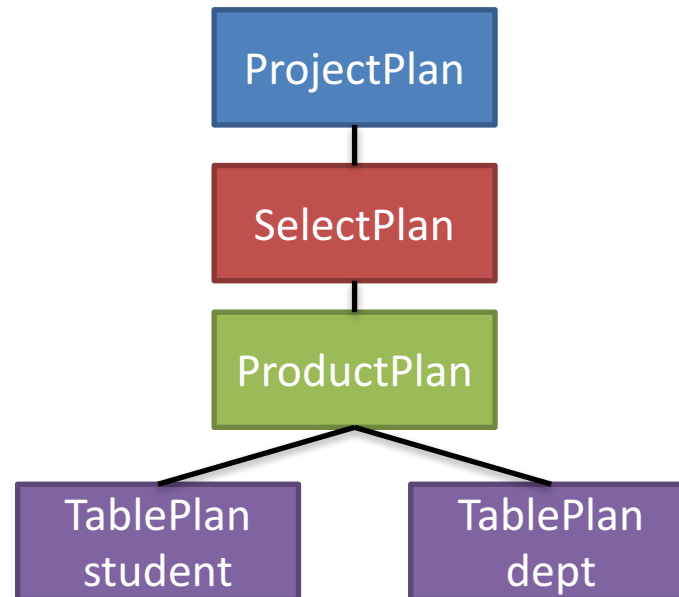



Why Does It Matter?

- A good scan tree can be faster than a bad one for orders of magnitude
- Consider the product scan at middle
 - Let $R(\text{student})=10000$, $B(\text{student})=1000$, $B(\text{dept})= 500$, and $\text{selectivity}(\text{s-id}=5\&\text{major-id}=4)=0.01$
 - Each block access requires 10ms
- Left: $(1000+10000*500)*10\text{ms} = 13.9 \text{ hours}$
- Right: $(1000+10000*0.01*500)*10\text{ms} = 8.4 \text{ mins}$
- We need a way to estimate the cost of a scan tree ***without actual scanning***
 - As we just did above

The Plan Interface

- A cost estimator for a *partial query*
- Each plan instance corresponds to an operator in relational algebra
 - Also to a subtree




Using a Query Plan

```
VanillaDb.init("studentdb");  
Transaction tx = VanillaDb.txMgr().transaction(  
    Connection.TRANSACTION_SERIALIZABLE, true);
```

```
Plan pb = new TablePlan("b", tx);  
Plan pu = new TablePlan("u", tx);  
Plan pp = new ProductPlan(pb, pu);  
Predicate pred = new Predicate(...);  
Plan sp = new SelectPlan(pp, pred);
```

```
sp.blockAccessed(); // estimate #blocks accessed
```

```
// open corresponding scan only if sp has low cost  
Scan s = sp.open();  
s.beforeFirst();  
while (s.next())  
    s.getVal("bid");  
s.close();
```



Plan before Scan

- A plan (tree) is a blueprint for evaluating a query
- Estimates cost by accessing statistics metadata only
 - No actual I/Os
 - Memory access only, *very efficient*
- Once a good plan is decided, we then create a scan following the blueprint

How to Find a Good Plan Tree?

- The planner can create multiple trees first, and then pick the one having the lowest cost
- Determining the best plan tree for a SQL command is call *planning*

Outline

- Overview
- Parsing and Validating SQL commands
 - Syntax vs. Semantics
 - Lexer, parser, and SQL data
 - Predicates
 - Verifier
- Scans and plans
- Query planning
 - Deterministic planners

What does a DB do when SQL coming?

1. Parses the SQL command
2. Verifies the SQL command
- 3. *Finds a good plan* for the SQL command**
4. Executes the plan

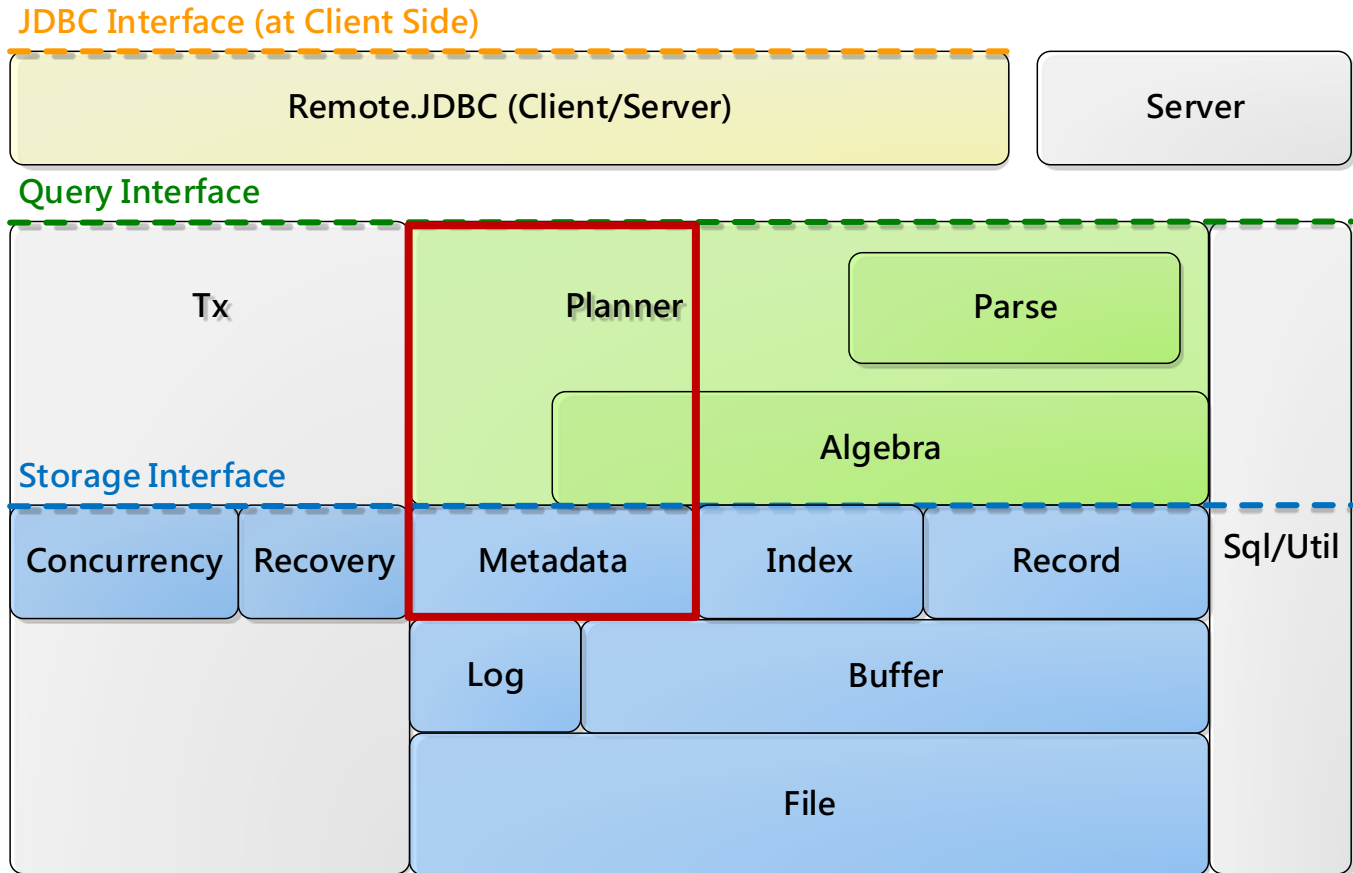
Planning

- Input:
 - SQL data
- Output:
 - A good plan tree
- Done by the *planner*
- How?

Query Optimization

Where Are We?

VanillaCore



Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

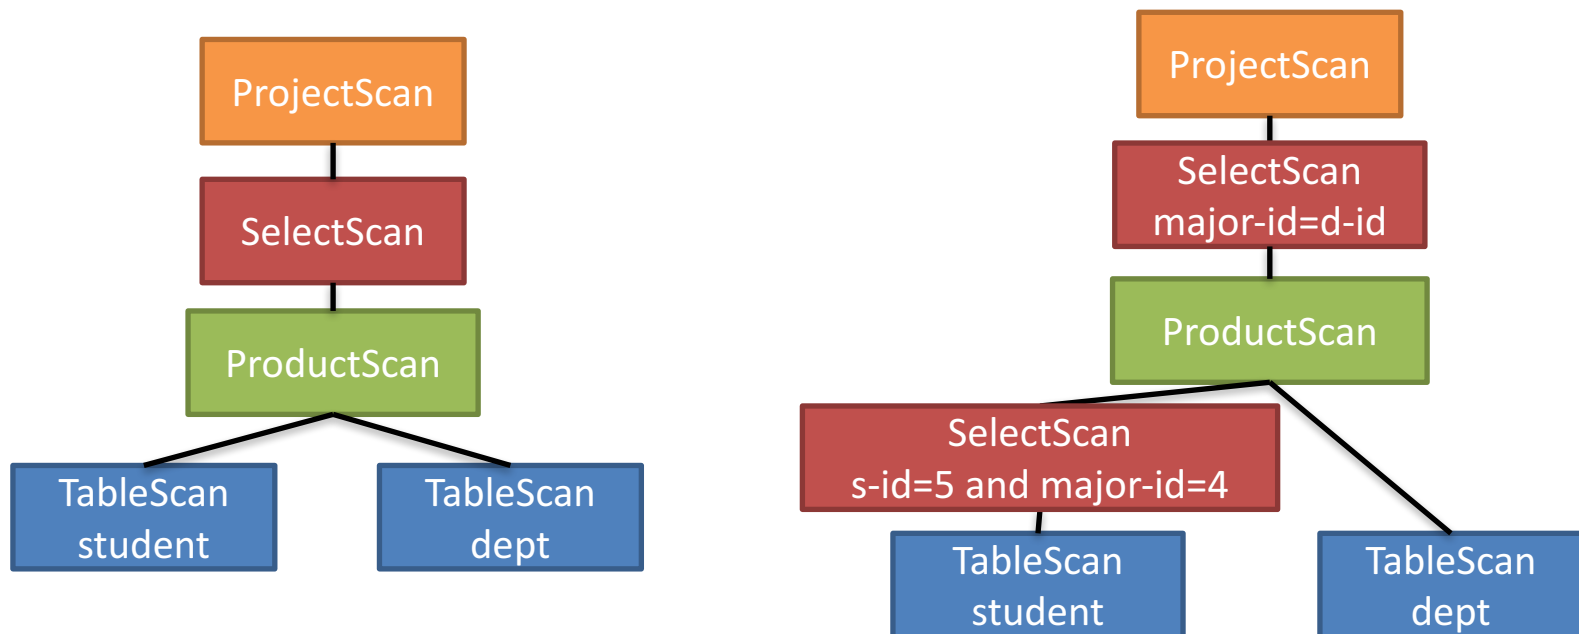
Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

SQL and Relational Algebra

- A SQL command can be expressed as multiple trees in relational algebra

```
SELECT sname FROM student, dept  
WHERE major-id = d-id AND s-id = 5 AND major-id = 4;
```



Query Optimization

- A good scan tree can be faster than a bad one for orders of magnitude
- Query optimizer (planner):
 1. Generate candidate plan trees
 2. Estimate cost of each corresponding scan tree (not opened yet)
 3. Pick and open the “best” one to execute query

Query Optimization

- A good scan tree can be faster than a bad one for orders of magnitude
- Query optimizer (planner):
 1. Generate candidate plan trees
 2. Estimate cost of each corresponding scan tree (not opened yet)
 3. Pick and open the “best” one to execute query

Outline

- Overview
- **Cost Estimation**
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Push Select Down
 - Join Order Problem
 - Heuristic Planner in VanillaCore
- Selinger-Style Query Optimizer

Metric for Cost

- Cost of query is generally measured as total elapsed time for answering query
 - Typically, I/O delay is the most important factor

Cost Estimation

- For each plan p , we estimate $B(p)$
 - #blocks accessed by the corresponding scan
- Usually, estimating $B(p)$ requires more knowledge:
 - $R(p)$: #records output
 - $V(p, f)$: #distinct values for field f
 - Search cost (#blocks) of index, if used

Estimating $B(p)$

p	$B(p)$
TablePlan	Actual #blocks cached by StatMgr (via periodic table scanning)
ProjectPlan(c)	$B(c)$
SelectPlan(c)	$B(c)$
IndexSelectPlan(t)	$\text{IndexSearchCost}(R(t), R(p)) + R(p)$
ProductPlan(c1, c2)	$B(c1) + (R(c1) * B(c2))$
IndexJoinPlan(c1, t2)	$B(c1) + (R(c1) * \text{IndexSearchCost}(R(t2), 1)) + R(p)$

Estimating $R(p)$ and Index Search Cost


- Index search cost:
 - `HashIndex.searchCost()`
 - `BTreeIndex.searchCost()`
- Estimating $R(p)$ is called ***cardinality estimation***

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Naïve Approach

- Uniform assumption
 - All values in field appear with the same probability



- Few statistics are enough:

R(c)	#records in child plan c
V(c, f)	#distinct values in field f in c
Max(c, f)	Max value in field f in c
Min(c, f)	Min value in field f in c

$$p = \text{Select}(c, f=x)$$

- $R(p)$?

$R(c)$	#records in child plan c
$V(c, f)$	#distinct values in field f in c
$\text{Max}(c, f)$	Max value in field f in c
$\text{Min}(c, f)$	Min value in field f in c

- $\text{Selectivity}(f=x): \frac{1}{V(c, f)}$
- $R(p): \text{Selectivity}(f=x) * R(c)$

$$p = \text{Select}(c, f > x)$$

- $R(p)$?

$R(c)$	#records in child plan c
$V(c, f)$	#distinct values in field f in c
$\text{Max}(c, f)$	Max value in field f in c
$\text{Min}(c, f)$	Min value in field f in c


- $\text{Selectivity}(f > x): \frac{\text{Max}(c, f) - x}{\text{Max}(c, f) - \text{Min}(c, f)}$
- $R(p): \text{Selectivity}(f > x) * R(c)$

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - **Histogram-based Estimation**
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer


Naïve Estimation is Inaccurate

- In the real world, values in a field are seldom uniform distributed
- $p = \text{Select}(c, f=14)$
- Estimated $R(p) = \frac{1}{15} * R(c) = 3$
- Actually, $R(p) = 9$



Histogram


- Approximates value distribution in every field
- Partitions field values into a set of *buckets*



- More #buckets, more accurate approximation
 - Tradeoff between accurate and storage cost

Buckets

- Each bucket collects statistics of a value range
 - Assuming *uniform distribution* within the bucket



- $R(p, f, b)$: #records
- $V(p, f, b)$: #distinct values
- $\text{Range}(p, f, b)$: value range

Cardinality Estimation

- Not matter what p is, we have

$$R(p) = \sum_{b \in p.hist.buckets(f)} R(p, f, b)$$


for any f

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - **Types of Histograms**
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer


“Raw” Histogram of a Table

- Data structure that approximates value distribution
- Partitions field values into a set of **buckets**
- Each bucket collects statistics of a value range
 - Assuming **uniform distribution** of records within the bucket
- Given raw values and #buckets, how to decide bucket ranges?



Equi-Width Histogram


- Partition strategy: all buckets have the same range
- $|\text{Range}(b)| = \frac{\text{Max}(p,f) - \text{Min}(p,f) + 1}{\#Buckets}$



- Problem: some buckets may be wasted

Equi-Depth Histogram


- Partition strategy: all buckets have the same #recs
- $\text{Depth} = \frac{R(p)}{\#Buckets}$



- Problem: records/values in a bucket may **not** be uniformly distributed

Max-Diff Histogram


- Partition strategy: split buckets at values with max. diff in #rec ($\text{MaxDiff}(F)$) or area ($\text{MaxDiff}(A)$):
 1. #recs: uniform #records in each bucket





Max-Diff Histogram

- Partition strategy: split buckets at values with max. diff in #rec (MaxDiff(F)) or area (MaxDiff(A)):

1. #recs: uniform #records in each bucket



2. Area: uniform #records *and values* in each bucket



Histogram in VanillaCore

- “Raw” histograms are statistics metadata
 - `org.vanilladb.core.storage.metadata.statistics`
- Accessed (by TablePlan) via `StatMgr.getTableStatInfo()`

Histogram
<pre>+ Histogram() + Histogram(fldnames : Set<String>) ~ Histogram(dists : Map<String, Collection<Bucket>>) + Histogram(hist : Histogram) + fields() : Set<String> + buckets(fldname : String) : Collection<Bucket> + addField(fldname : String) + addBucket(fldname : String, bkt : Bucket) + setBuckets(fldname : String, bkts : Collection<Bucket>) + recordsOutput() : double + distinctValues(fldname : String) : double + toString() : String + toString(int) : String</pre>

Bucket
<pre>+ Bucket(valrange : ConstantRange, freq : double, distvals : double) + Bucket(valrange : ConstantRange, freq : double, distvals : double, pcts : Percentiles) + valueRange() : ConstantRange + frequency() : double + distinctValues() : double + distinctValues(range : ConstantRange) : double + valuePercentiles() : Percentiles + toString() : String + toString(int) : String</pre>

Building Histogram (1/2)

- When system starts up:
- StatMgr:
 - Scans table and calls
`SampledHistogramBuilder.sample()`
 - When done, calls
`SampledHistogramBuilder.newMaxDiffHistogram()`
- Histogram types:
 - `MaxDiff(A)` : when field value is numeric
 - `MaxDiff(F)` : otherwise

Building Histogram (2/2)


- At runtime:
- StatMgr tacks #recs updated for each table
 - QueryPlanner calls StatMgr.countRecordUpdates() after executing modify/insert/delete queries
- Rebuilds histogram *in background* when StatMgr.getTableStatInfo() is called
 - If #recs updated > threshold (e.g., 100)
- StatisticsRefreshTask:
 - Scans table and calls SampledHistogramBuilder.sample()
 - When done, calls SampledHistogramBuilder.newMaxDiffHistogram()

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Query Optimization

- Query optimizer:
 1. Generate candidate plan trees
 2. Estimate cost of each corresponding scan tree
 3. Pick and open the “best” one to execute query

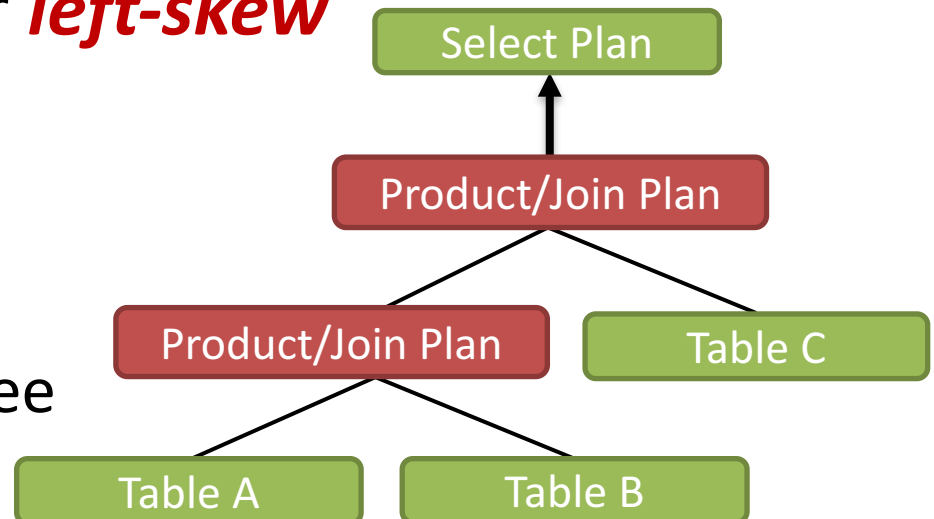


In Reality...

- Generating all candidate plan trees are too costly
 - #trees with n products/joins = Catalan number:

$$\frac{1}{n+1} \binom{2n}{n}$$

- Compromise: consider *left-skew* candidate trees only
- Query planner's goal
 - Avoiding bad trees
 - Not finding the best tree



Why Left-Skew Trees Only?

- Tend to be better than plans of other shapes
- Because many join algorithms scan right child c_2 multiple times
- Normally, we don't want c_2 to be a complex subtree

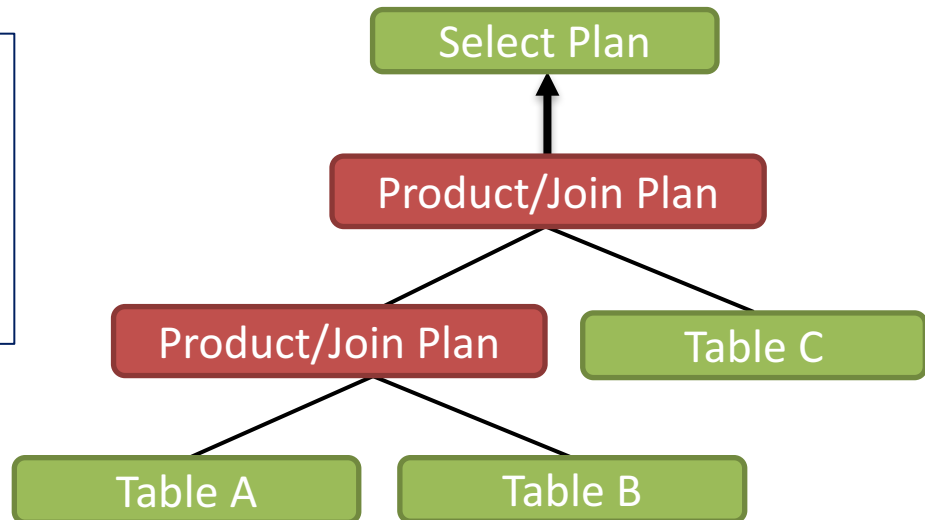
BasicQueryPlanner

```
public Plan createPlan(QueryData data, Transaction tx) {
    // Step 1: Create a plan for each mentioned table or view
    List<Plan> plans = new ArrayList<Plan>();
    for (String tblname : data.tables()) {
        String viewdef = VanillaDb.catalogMgr().getViewDef(tblname, tx);
        if (viewdef != null)
            plans.add(VanillaDb.newPlanner().createQueryPlan(viewdef, tx));
        else
            plans.add(new TablePlan(tblname, tx));
    }
    // Step 2: Create the product of all table plans
    Plan p = plans.remove(0);
    for (Plan nextplan : plans)
        p = new ProductPlan(p, nextplan);
    // Step 3: Add a selection plan for the predicate
    p = new SelectPlan(p, data.pred());
    // Step 4: Add a group-by plan if specified
    if (data.groupFields() != null) {
        p = new GroupByPlan(p, data.groupFields(), data.aggregationFn(), tx);
    }
    // Step 5: Project onto the specified fields
    p = new ProjectPlan(p, data.projectFields());
    // Step 6: Add a sort plan if specified
    if (data.sortFields() != null)
        p = new SortPlan(p, data.sortFields(), data.sortDirections(), tx);
    // Step 7: Add a explain plan if the query is explain statement
    if (data.isExplain())
        p = new ExplainPlan(p);
    return p;
}
```

- Product/join order follows what's written in SQL

Cost & Bottlenecks

```
SELECT  A.c1, B.c2, C.c3
FROM    A, B, C
WHERE   A.aid = C.aid
AND     B.bid = C.bid
AND     A.c2 = xxx
```

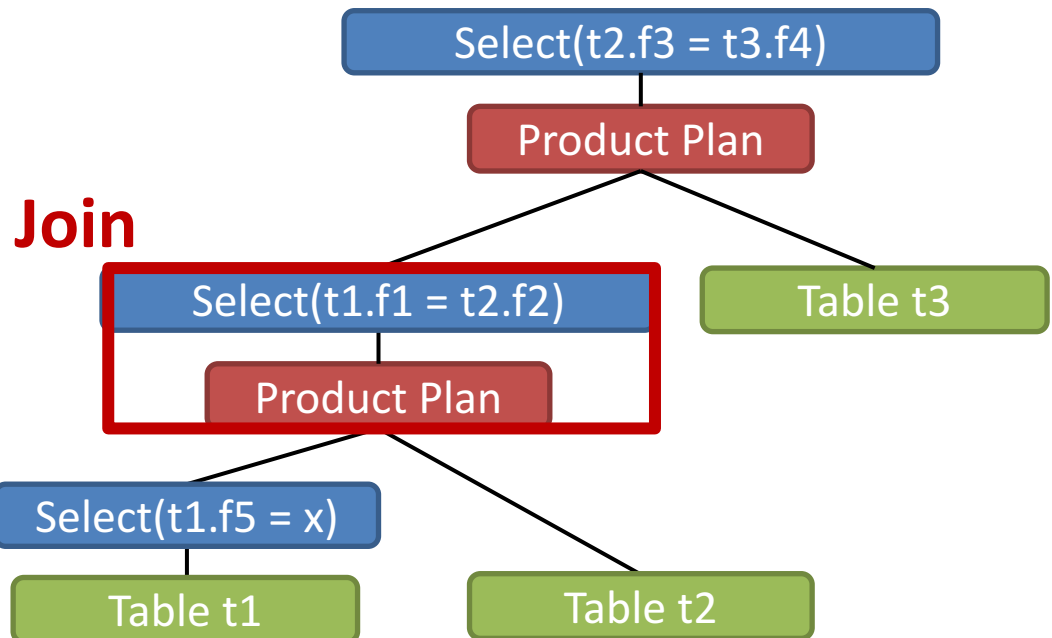


- B(root) dominated by **#recs** of product/join ops
 - $B(\text{Product}(c1, c2)) = B(c1) + (R(c1) * B(c2))$
 - $B(\text{IndexJoin}(c1, c2)) = B(c1) + (R(c1) * \text{SearchCost}(...)) + \dots$
 - $B(\text{Select}(c)) = B(c)$

Pushing Select Ops Down


- Execute Select ops as early as possible
- $\downarrow R(c1)$ and $\downarrow R(c2)$ of each product/join op

```
SELECT *  
FROM   t1, t2, t3  
WHERE  t1.f1 = t2.f2  
AND    t2.f3 = t3.f4  
AND    t1.f5 = x
```



Greedy Join Ordering

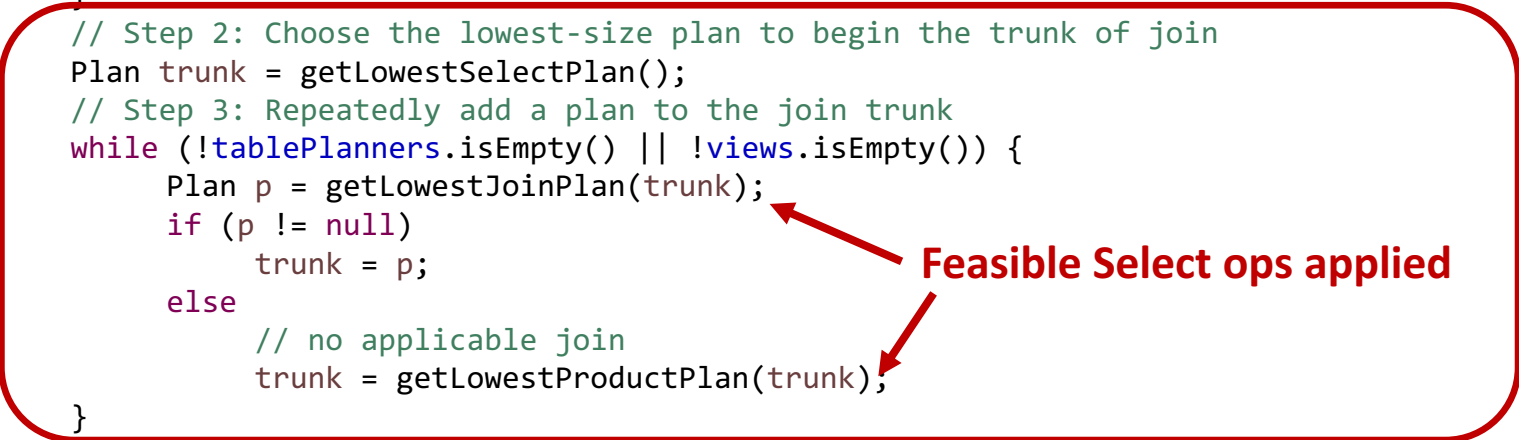
- $B(\text{root}) = B(p1) + (R(p1) * \dots) + \dots$
 - $\downarrow B(\text{root})$ implies $\downarrow(p1)$
- $B(p1) = B(c1) + (R(c1) * \dots) + \dots$
 - $\downarrow B(\text{root})$ also implies $\downarrow(c1)$
- ...
- $B(\text{root}) \propto R(p1) + R(c1) + \dots$



- ***Greedy Join ordering***: repeatedly add table to the “trunk” that result in lowest $R(\text{trunk})$

HeuristicPlanner in VanillaCore

```
public Plan createPlan(QueryData data, Transaction tx) {  
    // Step 1: Create a TablePlanner object for each mentioned table/view  
    int id = 0;  
    for (String tbl : data.tables()) {  
        String viewdef = VanillaDb.catalogMgr().getViewDef(tbl, tx);  
        if (viewdef != null)  
            views.add(VanillaDb.newPlanner().createQueryPlan(viewdef, tx));  
        else {  
            TablePlanner tp = new TablePlanner(tbl, data.pred(), tx, id);  
            tablePlanners.add(tp);  
        }  
        id += 1;  
    }  
    // Step 2: Choose the lowest-size plan to begin the trunk of join  
    Plan trunk = getLowestSelectPlan();  
    // Step 3: Repeatedly add a plan to the join trunk  
    while (!tablePlanners.isEmpty() || !views.isEmpty()) {  
        Plan p = getLowestJoinPlan(trunk);  
        if (p != null)  
            trunk = p;  
        else  
            // no applicable join  
            trunk = getLowestProductPlan(trunk);  
    }  
    // Step 4: Add a group by plan if specified  
    // Step 5: Project on the field names  
    // Step 6: Add a sort plan if specified  
    // Step 7: Add a explain plan if the query is explain statement  
}
```




Feasible Select ops applied

Outline

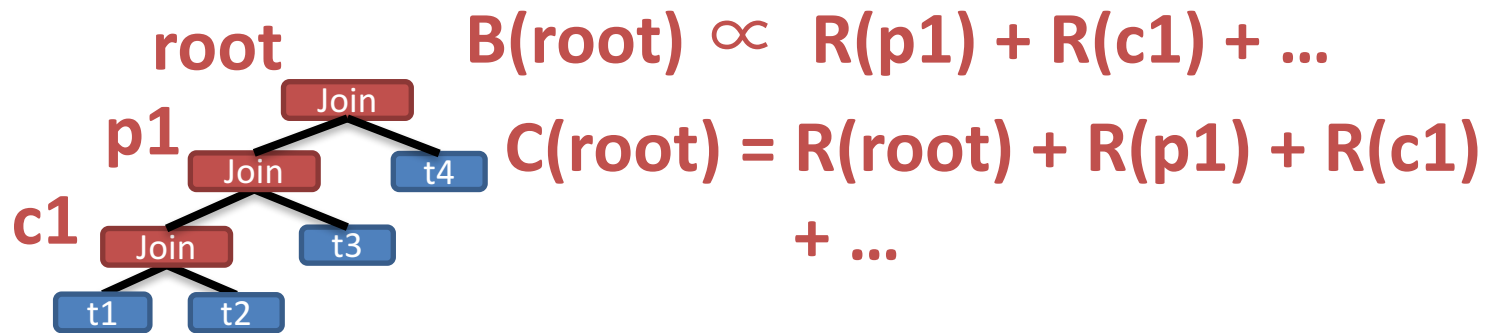
- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Why not HeuristicPlanner?



- Assumption: $\downarrow R(\text{c1})$ implies $\downarrow R(\text{p1})$
- May **not** be true: match rate matters
- Exhaustively searching the best join order?
 - Cost: **$O(n!)$** for n joins (e.g., $8! = 40320$)

Selinger-Style Optimizer



- Consider the best trees after 1, 2, 3, ... joins
- Observation: if $C(\mathbf{t3} \bowtie \mathbf{t1} \bowtie \mathbf{t2}) \leq C(\mathbf{t2} \bowtie \mathbf{t3} \bowtie \mathbf{t1})$, then $C(\mathbf{t3} \bowtie \mathbf{t1} \bowtie \mathbf{t2} \bowtie t4) \leq C(\mathbf{t2} \bowtie \mathbf{t3} \bowtie \mathbf{t1} \bowtie t4)$
- We can use *dynamic programming* to avoid repeating computations

Selinger Optimizer Example (1/3)

- Here are 3 relations to join: A, B, C
- Step 1:
 - compute the cost (R) of each relation's cheapest plan

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: {A} \bowtie {A, B}
 - Compare {B}A Cost: {B} \bowtie {B, A}

Same # of Record

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. $\{A, B\} =$
 - Compare $\{A\}B$ Cost: $\{A\} + \{A, B\}$
 - Compare $\{B\}A$ Cost: $\{B\} + \{B, A\}$

Choose the one with least cost

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. $\{A, B\} =$
 - Compare $\{A\}B$ Cost: $\{A\} + \{A, B\}$
 - Compare $\{B\}A$ Cost: $\{B\} + \{B, A\}$

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost

Selinger Optimizer Example (2/3)


- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: {A} + {A, B} = 159
 - Compare {B}A Cost: {B} + {B, A}

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20



2-Set	Best Plan	Cost

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: {A} + {A, B} = 159 
 - Compare {B}A Cost: {B} + {B, A} = 179

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: {A} + {A, B} = 159
 - Compare {B}A Cost: {B} + {B, A} = 179

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost
{A, B}	AB	159

Selinger Optimizer Example (2/3)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the record just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: {A} + (A, B) = 159
 - Compare {B}A Cost: {B} + (B, A) = 179

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare {A, B}C Cost: {A, B} + (A, B, C)
 - Compare {B, C}A Cost: {B, C} + (B, C, A)
 - Compare {C, A}B Cost: {C, A} + (C, A, B)

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77

3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare {A, B}C Cost: {A, B} + (A, B, C) = 259
 - Compare {B, C}A Cost: {B, C} + (B, C, A)
 - Compare {C, A}B Cost: {C, A} + (C, A, B)

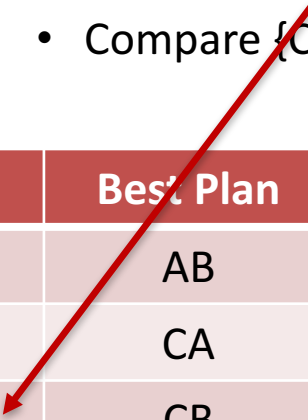
2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77

3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare {A, B}C Cost: {A, B} + (A, B, C) = 259
 - Compare {B, C}A Cost: {B, C} + (B, C, A)
 - Compare {C, A}B Cost: {C, A} + (C, A, B)

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77



3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare {A, B}C Cost: {A, B} + (A, B, C) = 259
 - Compare {B, C}A Cost: {B, C} + (B, C, A) = 111
 - Compare {C, A}B Cost: {C, A} + (C, A, B)

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77

3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)


- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. $\{A, B, C\} =$
 - Compare $\{A, B\}C$ Cost: $\{A, B\} + (A, B, C) = 259$
 - Compare $\{B, C\}A$ Cost: $\{B, C\} + (B, C, A) = 111$
 - Compare $\{C, A\}B$ Cost: $\{C, A\} + (C, A, B)$

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77



3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. $\{A, B, C\} =$
 - Compare $\{A, B\}C$ Cost: $\{A, B\} + (A, B, C) = 259$
 - Compare $\{B, C\}A$ Cost: $\{B, C\} + (B, C, A) = 111$
 - Compare $\{C, A\}B$ Cost: $\{C, A\} + (C, A, B) = 100$ 

2-Set	Best Plan	Cost
{A, B}	AB	159
{A, C}	CA	98
{B, C}	CB	77

3-Set	Best Plan	Cost

Selinger Optimizer Example (3/3)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. $\{A, B, C\} =$
 - Compare $\{A, B\}C$ Cost: $\{A, B\} + (A, B, C) = 259$
 - Compare $\{B, C\}A$ Cost: $\{B, C\} + (B, C, A) = 111$
 - Compare $\{C, A\}B$ Cost: $\{C, A\} + (C, A, B) = 100$

2-Set	Best Plan	Cost
$\{A, B\}$	AB	159
$\{A, C\}$	CA	98
$\{B, C\}$	CB	77

3-Set	Best Plan	Cost
$\{A, B, C\}$	CBA	100

Compare with Heuristic Planner

- Heuristic Planner
 - Greedy Search

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

- Selinger Optimizer
 - All combination

3-Set	Best Plan	Cost
{A, B, C}	CBA	100

```

private Plan getAllCombination(Plan viewTrunk) {
    long finalKey = 0;

    // for layer = 1, use select down strategy to construct
    for (TablePlanner tp: tablePlanners) {
        Plan bestPlan = null;
        if (viewTrunk != null) {
            bestPlan = tp.makeJoinPlan(viewTrunk);
            if (bestPlan == null)
                bestPlan = tp.makeProductPlan(viewTrunk);
        }
        else
            bestPlan = tp.makeSelectPlan();

        AccessPath ap = new AccessPath(tp, bestPlan);
        lookupTbl.put(ap.getAPIId(), ap);

        // compute final access path id
        finalKey += ap.getAPIId();
    }

    .
    .
    .

}

```

```

// construct all combination layer by layer
for (int layer = 2; layer <= tablePlanners.size(); layer++) {
    Set<Long> keySet = new HashSet<Long>(lookupTbl.keySet());

    for (TablePlanner rightOne: tablePlanners) {
        for (Long key: keySet) {
            AccessPath leftTrunk = lookupTbl.get(key);

            // cannot join with table which (layer-1) combination already included
            if (leftTrunk.isUsed(rightOne.getId()))
                continue;

            // do join
            Plan bestPlan = rightOne.makeJoinPlan(leftTrunk.getPlan());
            if (bestPlan == null)
                bestPlan = rightOne.makeProductPlan(leftTrunk.getPlan());

            AccessPath candidate = new AccessPath(leftTrunk, rightOne, bestPlan);
            AccessPath ap = lookupTbl.get(candidate.getAPIId());

            // there is no access path contains this combination
            if (ap == null) {
                lookupTbl.put(candidate.getAPIId(), candidate);
            }
            // check whether new access path is better than previous
            else {
                if (candidate.getCost() < ap.getCost())
                    lookupTbl.put(candidate.getAPIId(), candidate);
            }
        }
    }

    // remove the elements belong to layer-1
    // because in the next layer we only need this layer's combination
    for (Long key: keySet)
        lookupTbl.remove(key);
}

return lookupTbl.get(finalKey).getPlan();

```

- Iterate all table planners to join with all existing (layer-1) combination to construct this layer

```

public class AccessPath {
    private Plan p;
    private AccessPathId apId;
    private long cost = 0;
    private ArrayList<Integer> tblUsed = new ArrayList<Integer>();

```

```

    public class AccessPathId {
        long id;

        AccessPathId(TablePlanner tp) {
            this.id = (long) Math.pow(2, tp.getId());
        }

        AccessPathId(AccessPath ap, TablePlanner tp) {
            this.id = ap.getAPId() + (long) Math.pow(2, tp.getId());
        }

        public long getID() {
            return id;
        }
    }
}

```

- Using **sum of pow(2, tp.id)** to represent the combination of tables in this access path
- Using **pow(2, tp.id)** to avoid problems with different combinations but with the same apID
- Then we can use apID as the key of the lookup table

```

    public AccessPath (TablePlanner newTp, Plan p) {
        this.p = p;
        this.tblUsed.add(newTp.getId());
        this.apId = new AccessPathId(newTp);
        this.cost = p.recordsOutput();
    }

    public AccessPath (AccessPath preAp, TablePlanner newTp, Plan p) {
        this.p = p;
        this.tblUsed.addAll(preAp.getTblUsed());
        this.tblUsed.add(newTp.getId());
        this.apId = new AccessPathId(preAp, newTp);

        // approximate cost = previous cost + new cost
        this.cost = preAp.getCost() + p.recordsOutput();
    }
}

```

```

public class AccessPath {
    private Plan p;
    private AccessPathId apId;
    private long cost = 0;
    private ArrayList<Integer> tblUsed = new ArrayList<Integer>();

    public class AccessPathId {
        long id;

        AccessPathId(TablePlanner tp) {
            this.id = (long) Math.pow(2, tp.getId());
        }

        AccessPathId(AccessPath ap, TablePlanner tp) {
            this.id = ap.getAPId() + (long) Math.pow(2, tp.getId());
        }
        public long getID() {
            return id;
        }
    }

    public AccessPath (TablePlanner newTp, Plan p) {
        this.p = p;
        this.tblUsed.add(newTp.getId());
        this.apId = new AccessPathId(newTp);
        this.cost = p.recordsOutput();
    }
    public AccessPath (AccessPath preAp, TablePlanner newTp, Plan p) {
        this.p = p;
        this.tblUsed.addAll(preAp.getTblUsed());
        this.tblUsed.add(newTp.getId());
        this.apId = new AccessPathId(preAp, newTp);

        // approximate cost = previous cost + new cost
        this.cost = preAp.getCost() + p.recordsOutput();
    }
}

```

- Using **sum of pow(2, tp.id)** to represent the combination of tables in this access path
- Using **pow(2, tp.id)** to avoid problems with different combinations but with the same apID
- Then we can use apID as the key of the lookup table
- Approximate B(root) using $R(p1) + R(c1) \dots$

Reference

- <https://db.inf.uni-tuebingen.de/staticfiles/teaching/ws1011/db2/db2-selectivity.pdf>
- <https://www.cise.ufl.edu/~adobra/approxqp/histograms2>
- <https://pdfs.semanticscholar.org/b024/0a44105fa0a0967d96d109aac9f021902ebb.pdf>