

# Introduction to VanillaDB

DB/AI Bootcamp

2018 Summer

DataLab, CS, NTHU

# RDBMS

- Definition: A ***Relational DBMS (RDBMS)*** is a DBMS that supports the relational model

# Outline

- Architecture of an RDBMS
- Query interfaces
  - SQL, JDBC, and native interface
- Storage interface
  - RecordFile and metadata

# Outline

- Architecture of an RDBMS
- Query interfaces
  - SQL, JDBC, and native interface
- Storage interface
  - RecordFile and metadata

# Architecture of an RDBMS

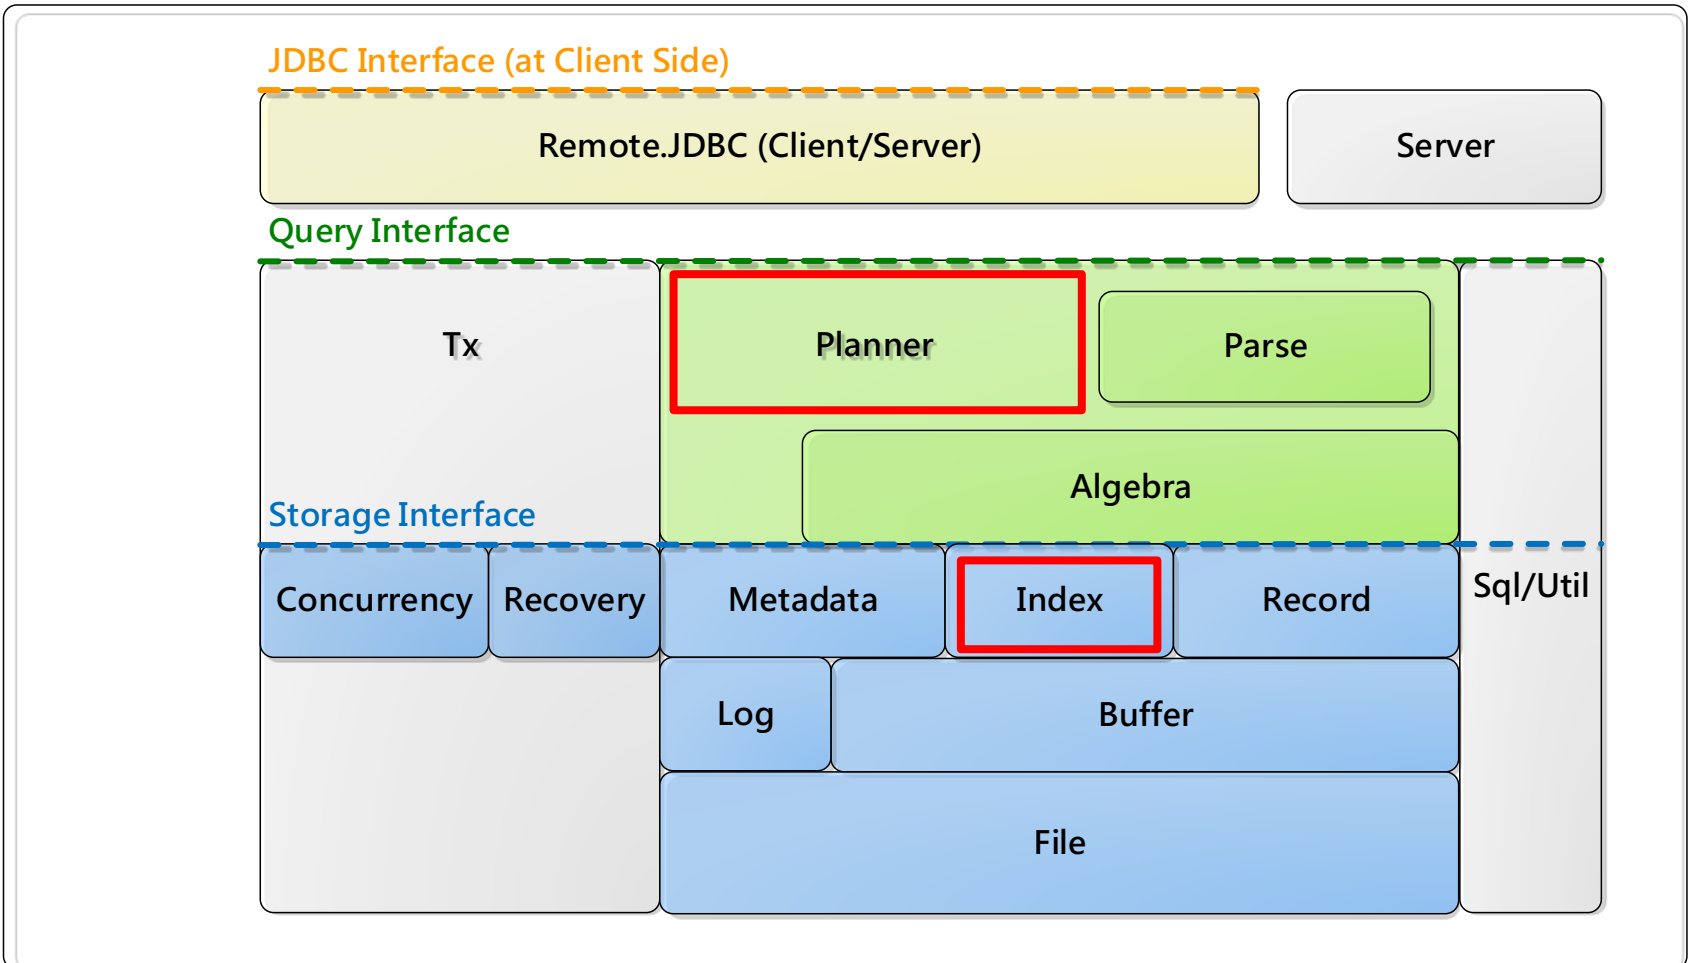
- Largely influenced by the IBM System R
  - Announced in 1974

# The VanillaDB Project

- VanillaCore
  - An RDBMS that runs on a single server
- VanillaComm
  - A communication infrastructure for distributed RDBMS

# Architecture of VanillaCore (1/2)

VanillaCore



# Architecture of VanillaCore (2/2)

- Interfaces:
  - SQL
  - JDBC
  - Native query interface
  - Storage interface (for file access)
- Key components:
  - Sever and infrastructures (jdbc, sql, tx, and utils)
  - Query engine
  - Storage engine



# Outline

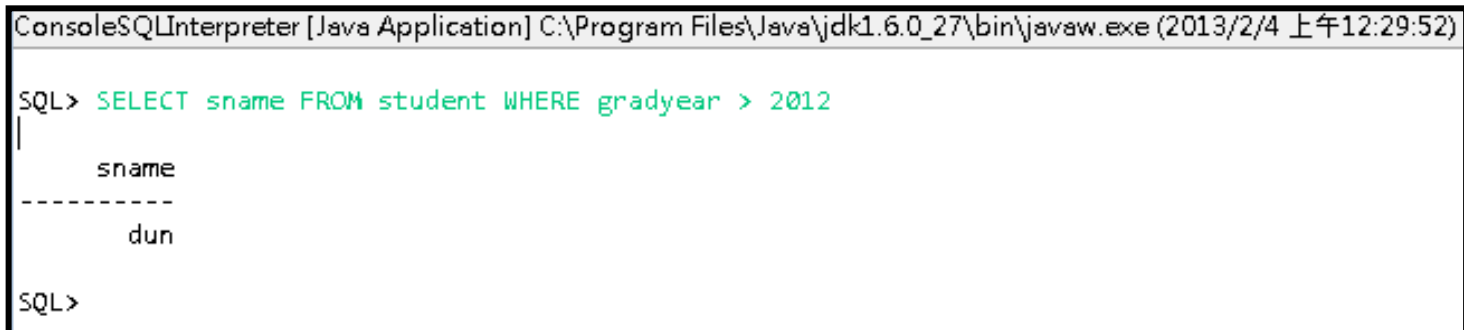
- Architecture of an RDBMS
- Query interfaces
  - SQL, JDBC, and Native
- Storage interface
  - RecordFile and metadata

# The SQL Interface

- **SQL** (Structured Query Language) is a standardized interface
  - SQL-92, SQL-99, and later versions

# Issuing SQL Commands

- Client-server mode:
  - Manual commands through `util.ConsoleSQLInterpreter`



The screenshot shows a Java application window titled "ConsoleSQLInterpreter [Java Application] C:\Program Files\Java\jdk1.6.0\_27\bin\javaw.exe (2013/2/4 上午12:29:52)". The window contains a text area with the following content:

```
SQL> SELECT sname FROM student WHERE gradyear > 2012
|
      sname
-----
      dun
SQL>
```

- Or in client programs through the **JDBC** interface
- Embedded mode:
  - Through the native query interface

# Supported SQL Commands (1/5)

- VanillaCore supports a tiny subset of SQL-92
  - DDL: CREATE <TABLE | VIEW | INDEX>
  - DML: SELECT, UPDATE, INSERT, DELETE
- Limitations:
  - Types: int, long, double, and varchar
  - Single SELECT-FROM-WHERE block
    - No \* in SELECT clause, no AS in FROM, no null value, no explicit JOIN or OUTER JOIN, only AND in WHERE, no parentheses, no computed value
  - Arithmetic expression only in UPDATE
  - No query in INSERT

# Supported SQL Commands (2/5)

```
<Field>      := IdTok
<Constant>   := StrTok | NumericTok
<Expression> := <Field> | <Constant>
<BinaryArithmeticExpression> :=
    ADD(<Expression>, <Expression>) |
    SUB(<Expression>, <Expression>) |
    MUL(<Expression>, <Expression>) |
    DIV(<Expression>, <Expression>)
<Term>       := <Expression> = <Expression> |
    <Expression> > <Expression> |
    <Expression> >= <Expression> |
    <Expression> < <Expression> |
    <Expression> <= <Expression>
<Predicate>  := <Term> [ AND <Predicate> ]
```

# Supported SQL Commands (3/5)

**<Query>                := SELECT <ProjectSet> FROM <TableSet>  
                         [ WHERE <Predicate> ] [ GROUP BY <IdSet> ]  
                         [ ORDER BY <SortList> [ DESC | ASC ] ]**

**<IdSet>                := <Field> [ , <IdSet> ]**

**<TableSet>            := IdTok [ , <TableSet> ]**

**<AggFn>                := AVG(<Field>) | COUNT(<Field>) |  
                         COUNT(DISTINCT <Field>) | MAX(<Field>) |  
                         MIN(<Field>) | SUM(<Field>)**

**<ProjectSet>        := <Field> | <AggFn> [ , <ProjectSet>]**

**<SortList>            := <Field> | <AggFn> [ , <SortList>]**

# Supported SQL Commands (4/5)

`<UpdateCmd> := <Insert> | <Delete> | <Modify> | <Create>`  
`<Create> := <CreateTable> | <CreateView> |`  
`<CreateIndex>`  
`<Insert> := INSERT INTO IdTok ( <FieldList> ) VALUES`  
`( <ConstantList> )`  
`<FieldList> := <Field> [ , <Field> ]`  
`<ConstantList>:= <Constant> [ , <Constant> ]`  
`<Delete> := DELETE FROM IdTok [ WHERE <Predicate> ]`  
`<Modify> := UPDATE IdTok SET <ModifyTermList>`  
`[ WHERE <Predicate> ]`

# Supported SQL Commands (5/5)

```

<ModifyExpression>    := <Expression> |
                        <BinaryArithmeticExpression>
<ModifyTermList>     := <Field> = <ModifyExpression>
                        [ , <ModifyTermList> ]

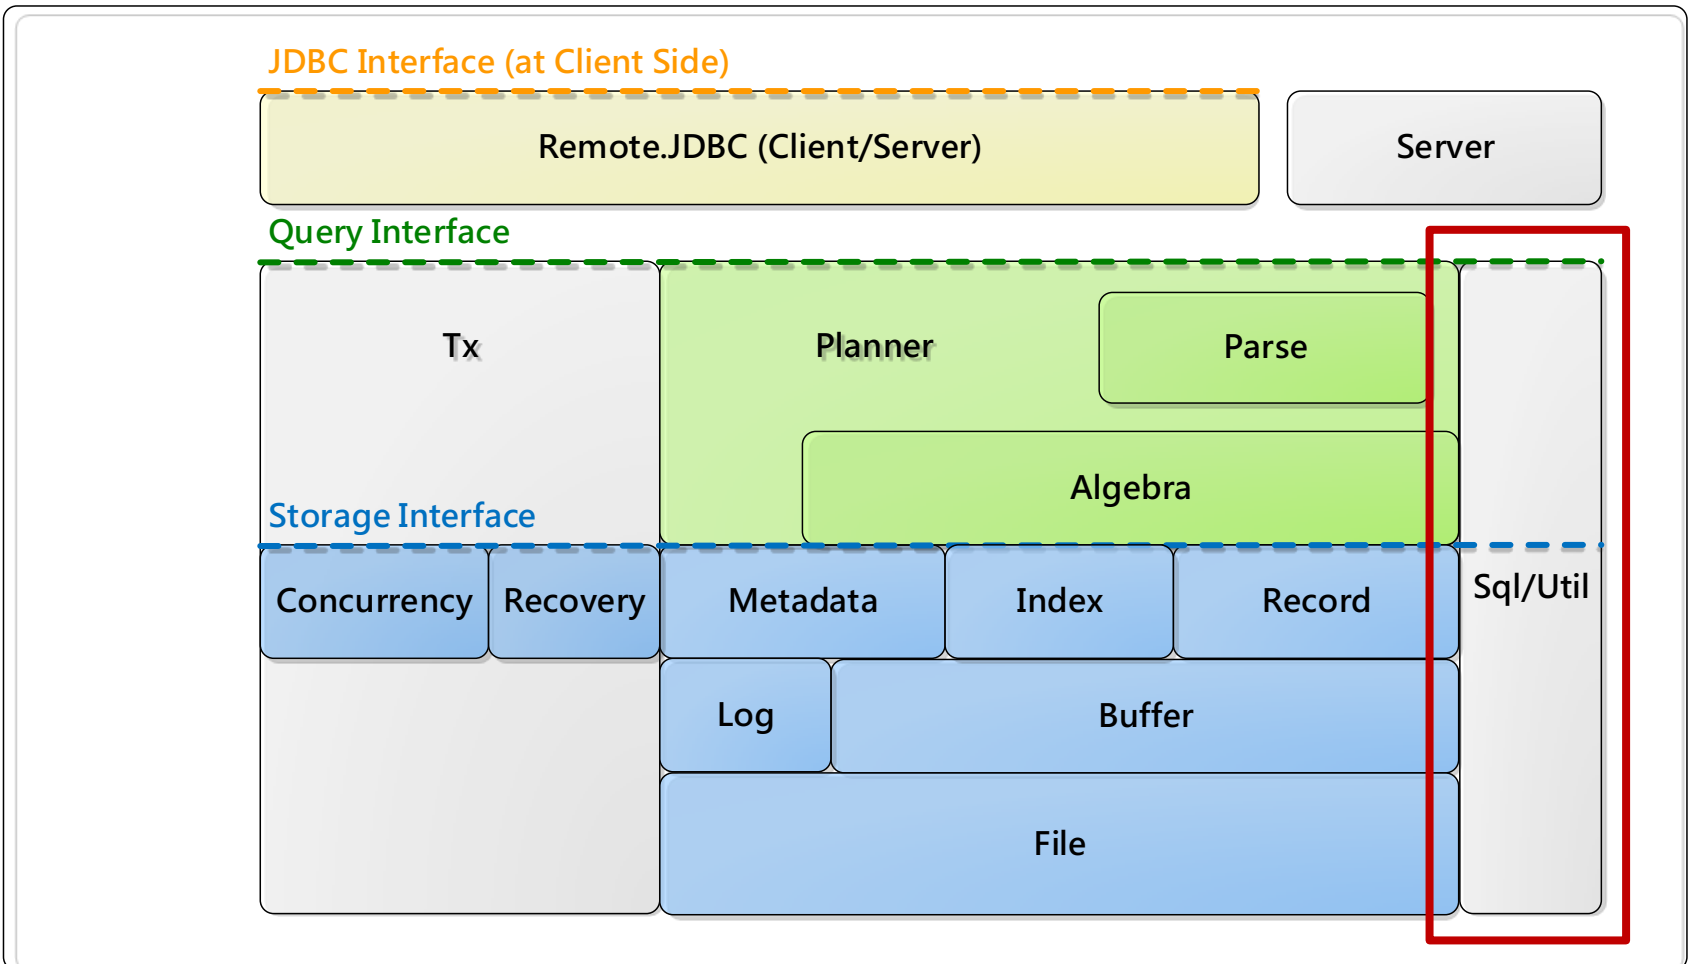
<CreateTable>         := CREATE TABLE IdTok ( <FieldDefs> )
<FieldDefs>          := <FieldDef> [ , <FieldDef> ]
<FieldDef>            := IdTok <TypeDef>
<TypeDef>             := INT | LONG | DOUBLE |
                        VARCHAR ( NumericTok )
<CreateView>          := CREATE VIEW IdTok AS <Query>
<CreateIndex>         := CREATE INDEX IdTok ON IdTok
                        ( <Field> )

```



# Architecture of VanillaCore

VanillaCore



# Utility Classes for SQL

- In `sql` package
- **Types:**
  - **Numeric:** `IntegerType`, `BigIntType`, and `DoubleType`
  - **String:** `VarcharType`
- **Constants:**
  - `IntegerConstant`, `BigIntConstant`, `DoubleConstant`, **and** `VarcharConstant`
- **For relations:**
  - `Schema`, `Record`
- **For commands:**
  - `Predicate`, `AggFn`

# Types

- Each `Type` impl. denotes a supported SQL type

<code>&lt;&lt;abstract&gt;&gt;</code> <code>Type</code>
<u><code>&lt;&lt;final&gt;&gt; + INTEGER : Type</code></u> <u><code>&lt;&lt;final&gt;&gt; + BIGINT : Type</code></u> <u><code>&lt;&lt;final&gt;&gt; + DOUBLE : Type</code></u> <u><code>&lt;&lt;final&gt;&gt; + VARCHAR : Type</code></u>
<u><code>+ VARCHAR(arg : int) : Type</code></u> <u><code>+ newInstance(sqlType : int) : Type</code></u> <u><code>+ newInstance(sqlType : int, arg : int) : Type</code></u> <code>&lt;&lt;abstract&gt;&gt; + getSqlType() : int</code> <code>&lt;&lt;abstract&gt;&gt; + getArgument() : int</code> <code>&lt;&lt;abstract&gt;&gt; + isFixedSize() : boolean</code> <code>&lt;&lt;abstract&gt;&gt; + isNumeric() : boolean</code> <code>&lt;&lt;abstract&gt;&gt; + maxSize() : int</code> <code>&lt;&lt;abstract&gt;&gt; + maxValue() : Constant</code> <code>&lt;&lt;abstract&gt;&gt; + minValue() : Constant</code>

<code>java.sql.Types</code>	<code>vanilladb.sql.Type</code>
INTEGER	IntegerType
BIGINT	BigIntType
DOUBLE	DoubleType
VARCHAR	VarcharType

# Constants

- Each `Constant` impl. denotes a value of a supported type
  - Immutable
  - Arithmetics with auto type-upgrade

<<abstract>> Constant
<u>+ newInstance(type : Type, val : byte[]) : Constant</u> <u>+ defaultInstance(type : Type) : Constant</u> <<abstract>> + getType() : Type <<abstract>> + asJavaVal() : Object <<abstract>> + asBytes() : byte[] <<abstract>> + size() : int <<abstract>> + castTo(type : Type) : Constant <<abstract>> + add(c : Constant) : Constant <<abstract>> + sub(c : Constant) : Constant <<abstract>> + mul(c : Constant) : Constant <<abstract>> + div(c : Constant) : Constant

vanilladb.sql.Type	Value type in Java
IntegerType	Integer
BigIntType	Long
DoubleType	Double
VarcharType	String

# Relations

blog\_pages

blog_id	url	created	author_id
33981	ms.com/...	2012/10/31	729
33982	apache.org/...	2012/11/15	4412

← Schema

← Record

# Schema & Record

Schema
+ Schema() + addField(fldname : String, type : Type) + add(fldname : String, sch : Schema) + addAll(sch : Schema) + fields() : SortedSet<String> + hasField(fldname : String) : boolean + type(fldname : String) : Type

- Contains the name and type of each field in a table

<<interface>> Record
+ getVal(fldName : String) : Constant

- A map from field names to constants

# Commands

- **Supporting WHERE: predicates in `sql.predicate` package**
  - `Expression`, `FieldExpression`, `ConstantExpression`, `BinaryArithmeticExpression`, `Term`, and `Predicate`
- **Supporting GROUP BY: aggregation functions in the `sql.aggfn` package**
  - `AggregationFn`, `AvgFn`, `CountFn`, `DistinctCountFn`, `MaxFn`, `MinFn` and `SumFn`

# Outline

- Architecture of an RDBMS
- Query interfaces
  - SQL, JDBC, and native interface
- Storage interface
  - RecordFile and metadata

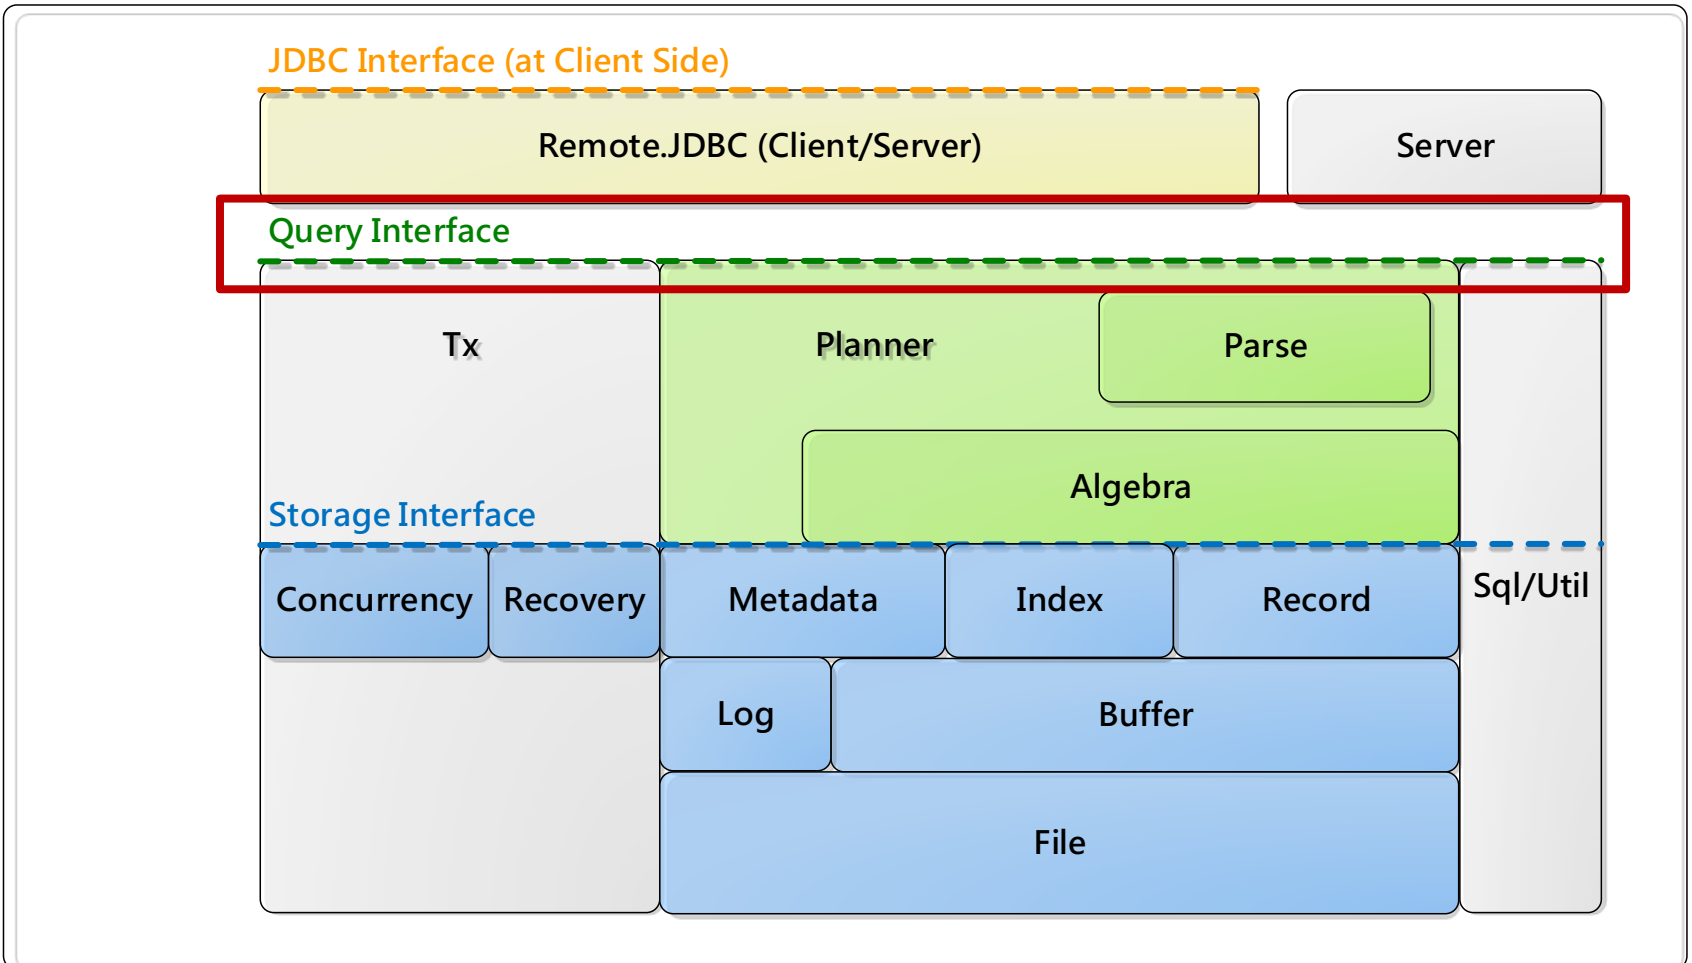


# Outline

- Architecture of an RDBMS
- Query interfaces
  - SQL, JDBC, and native interface
- Storage interface
  - RecordFile and metadata

# Architecture of VanillaCore (1/2)

VanillaCore



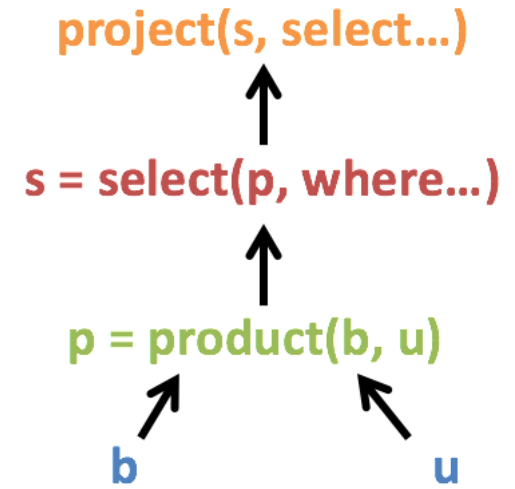
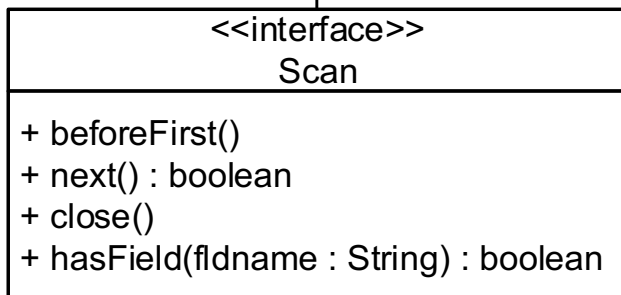
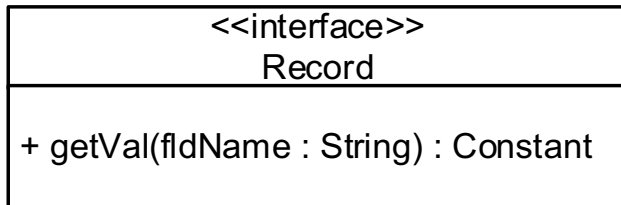
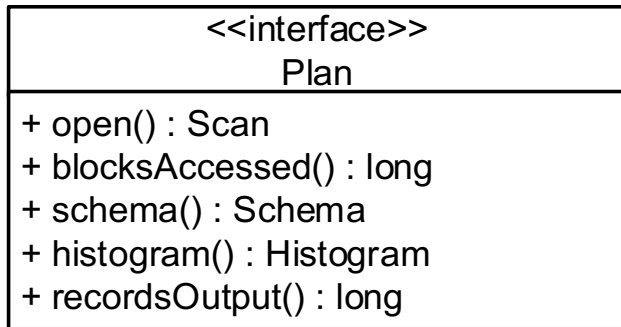
# API (1/2)

Planner
+ createQueryPlan(qry : String, tx : Transaction) : Plan + executeUpdate(cmd : String, tx : Transaction) : int

Transaction
+ <u>addStartListener</u> (l : TransactionLifeCycleListener) + Transaction(concurMgr : TransactionLifeCycleListener, recoveryMgr : TransactionLifeCycleListener, readOnly : boolean, txNum : long) + addLifeCycleListener(l : TransactionLifeCycleListener) + commit() + rollback() + recover() + endStatement() + getTransactionNumber() : long + isReadOnly() : boolean + concurrencyMgr() : ConcurrencyMgr + recoveryMgr() : RecoveryMgr

- All operations resulted from a planner are bound by the associated tx

# API (2/2)



- Corresponds to an operator in relational algebra
  - The root of a plan tree
  - For **cost estimation** only
  - `open()` propagates down to the tree
- Iterator of output records of a partial query
  - Actual data access

# References

- M.M. Astrahan et al., System R: relational approach to database management, *ACM Transactions on Database Systems*, Vol. 1, No. 2, 1976
- J. M. Hellerstein et al., Architecture of a database system, *Foundations and Trends in Databases*, Vol. 1, No. 2, 2007
- Edward Sciore, Chapters 8 & 20, *Database Design and Implementation*, 2008