

Transaction Management: Concurrency Control

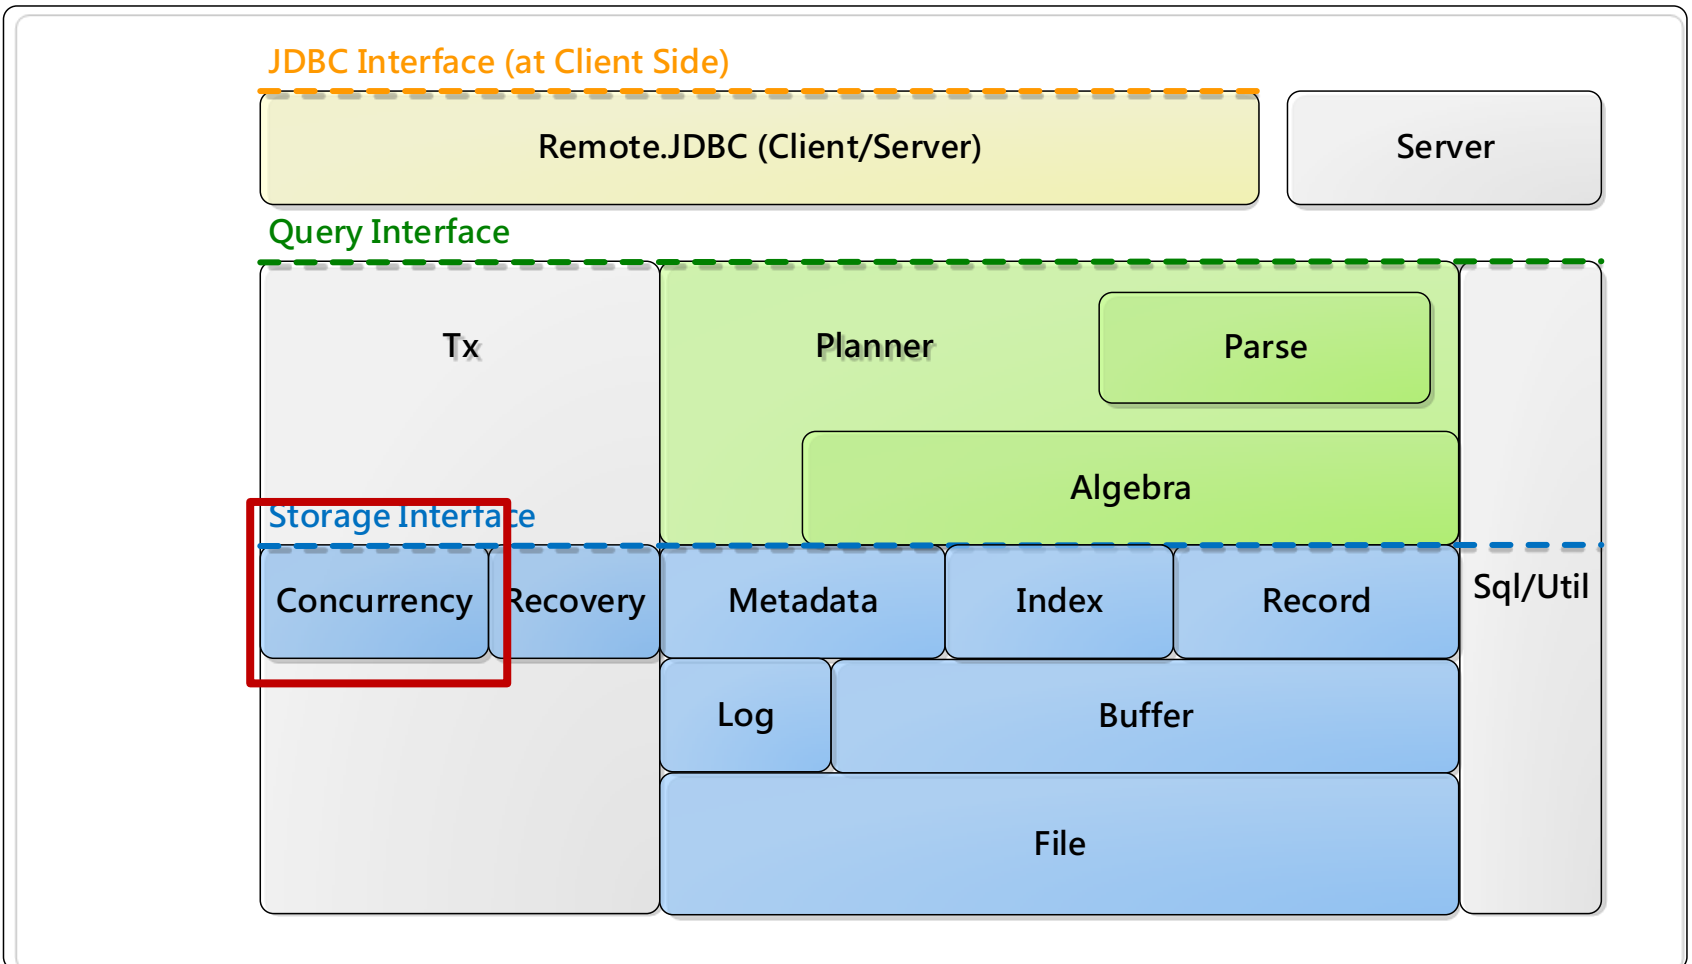
DB/AI Bootcamp

2018 Summer

Datalab, CS, NTHU

Concurrency Mgr

VanillaCore



Outline

- Schedules
- Anomalies
- Lock-based concurrency control

- 2PL and S2PL

Skipped.

- Deadlock

Check out NTHU CS 471000 if you are interested in.

- Granularity of locks

- Dynamic databases

- Phantom

- Isolation levels

- Meta-structures

- Concurrency manager in VanillaCore

Outline

- Schedules
- Anomalies
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

Outline

- Schedules
- Anomalies
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

Concurrency Manager

- Ensures *consistency* and *isolation*

Consistency

- ***Consistency***

- Txs will leave the database in a consistent state
- I.e., all integrity constraints (ICs) are meet
 - Primary and foreign key constrains
 - Non-null constrain
 - (Field) type constrain
 - ...
- Users are responsible for issuing “valid” txs

Isolation

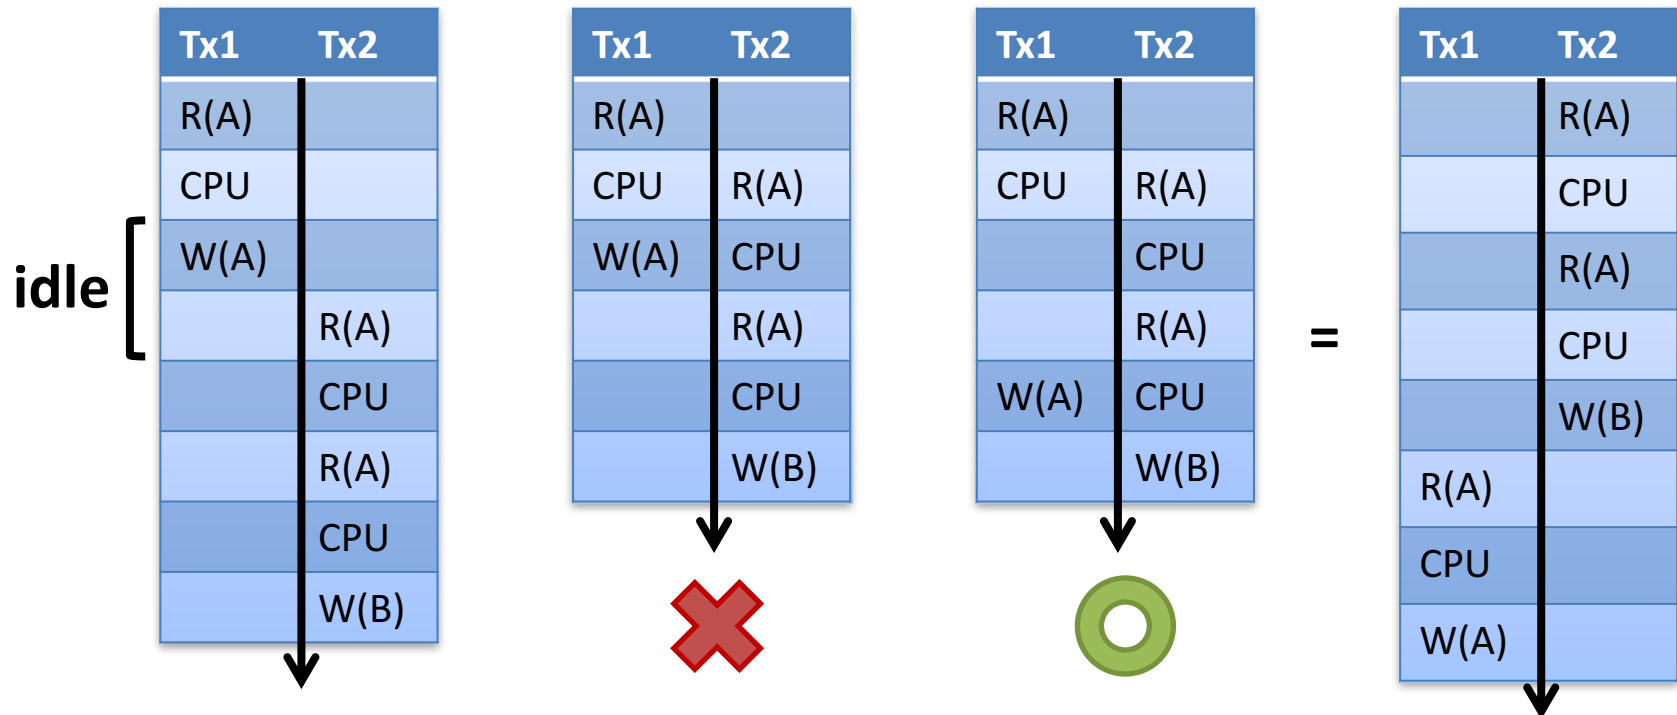
- ***Isolation***

- Interleaved execution of txs should have the net effect identical to executing tx in ***some*** serial order
- T_1 and T_2 are executed concurrently, isolation gives that the net effect to be equivalent to either
 - T_1 followed by T_2 or
 - T_2 followed by T_1
- The DBMS does ***not*** guarantee to result in ***which particular*** order

Why do we need to interleave txs?

Concurrent Txs

- Since I/O is slow, it is better to execute Tx1 and Tx2 concurrently to reduce CPU idle time



- The concurrent result should be the same as serial execution in *some* order
 - Better concurrency

Concurrent Txs

- Pros:
 - Increases throughput (via CPU and I/O pipelining)
 - Shortens response time for short txs
- But operations must be interleaved correctly

Transactions and Schedules

- Before executing T_1 and T_2 :
 - $A = 300, B = 400$

T1:	BEGIN	$A=A+100,$	$B=B-100$	END
T2:	BEGIN	$A=1.06*A,$	$B=1.06*B$	END

- Two possible execution results
 - T_1 followed by T_2
 - $A = 400, B = 300 \rightarrow A = 424, B = 318$
 - T_2 followed by T_1
 - $A = 318, B = 424 \rightarrow A = 418, B = 324$

Transactions and Schedules

- A *schedule* is a list of actions/operations from a set of transaction
- If the actions of different transactions are not interleaved, we call this schedule a *serial schedule*

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, \quad B = 1.06 * B$	

Transactions and Schedules

- Equivalent schedules
 - The effect of executing the first schedule is identical to the effect of executing the second schedule
- ***Serializable schedule***
 - A schedule that is equivalent to some serial execution of the transactions

Transactions and Schedules

- A possible interleaving schedule

T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A,$	$B=1.06*B$

– Result: $A = 424, B = 318$

– A serializable schedule

- Equivalent to T_1 followed by T_2

T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A,$	$B=1.06*B$

Transactions and Schedules

- How about this schedule?

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, B = 1.06 * B$	

- Result: **$A = 424, B = 324$**
- A non-serializable schedule
- Violates the isolation requirement

Goal

- Interleave operations while making sure the schedules are serializable
- How?

Outline

- Schedules
- **Anomalies**
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

Simplified Notation

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, B = 1.06 * B$	

- Can be simplified to:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	


- Here, we care about operations, not values

Anomalies

- Weird situations that would happen when interleaving operations
 - But not in serial schedules
- Mainly due to the *conflicting* operations


Conflict Operations

- Two operations on the same object are **conflict** if they are operated by different txs and at least one of these operations is a write



Types

- Write-read conflict
- Read-write conflict
- Write-write conflict



- Read-read conflict?
 - No anomaly

Anomalies due to Write-Read Conflict

- Reading uncommitted data

– *Dirty reads*

T1:	R(A), <u>W(A)</u> ,	R(B), W(B)
T2:	<u>R(A)</u> , W(A), R(B), W(B)	

- A *unrecoverable schedule*

T1:	R(A), W(A),	R(B), W(B), Abort
T2:	R(A), W(A), C	

– T1 cannot abort!


– *Cascading aborts* if T2 completes after T1 aborts

Anomalies due to Read-Write Conflict

- **Unrepeatable reads:**

- T_1 : if $(A > 0) A = A - 1$;
- T_2 : if $(A > 0) A = A - 1$;
- IC on A : cannot be negative

T1:	<u>R(A),</u>	<u>R(A),</u>	W(A), C
T2:		R(A),	<u>W(A),</u> C




Anomalies due to Write-Write Conflict

- ***Lost updates:***

- $T_1: A = A + 1; B = B * 10;$
- $T_2: A = A + 2; B = B * 5;$
- Start with $A=10, B=10$

T1:	<u>W(A),</u>	W(B), C
T2:	<u>W(A),</u>	W(B), C




Avoiding Anomalies

- Idea:
- Perform all conflicting actions between T1 and T2 *in the same order* (either T1's before T2's or T2's before T1's)
- I.e., to ensure *conflict serializability*


Conflict Equivalent

- If two operations are not conflict, we can **swap** them to generate an equivalent schedule
- Schedule 1 is **conflict equivalent** to schedule 2 and schedule 3


Schedule 1



Schedule 2




Schedule 3




Conflict Serializable

- By swapping non-conflict operations, we can transfer the schedule 1 into a serial schedule 4
- We say that schedule 1 is ***conflict serializable***


Schedule 3



Schedule 4



Ensuring Conflict Serializability is *Not Enough*



- Conflict serializable, but *not* recoverable

Avoiding Anomalies

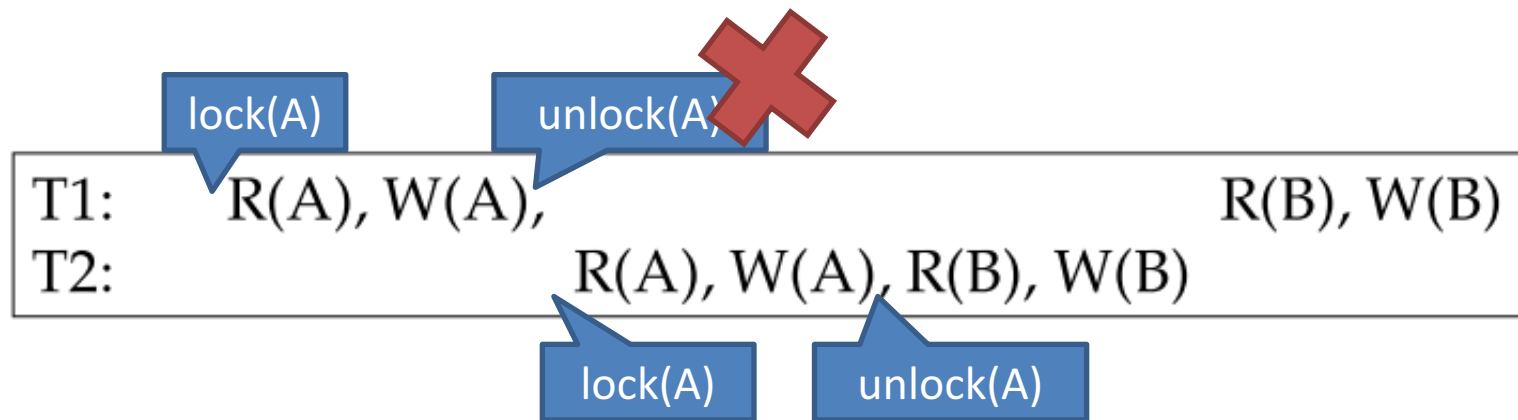
- We also need to ensure recoverable schedule
- Definition: A schedule is *recoverable* if each tx T commits only after all txs whose changes T reads, commit
- How?
 - Avoid cascading aborts
 - Disallow a tx from reading uncommitted changes from other txs

Outline

- Schedules
- Anomalies
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

Lock-Based Concurrency Control

- For isolation and consistency, a DBMS should only allow *serializable*, *recoverable* schedules
 - Uncommitted changes cannot be seen (no WR)
 - Ensure repeatable read (no RW)
 - Cannot overwrite uncommitted change (no WW)
- A *lock* for each data item seems to be a good solution



Lock \neq latch

- Lock: long-term, tx-level
- Latch: short-term, ds/alg-level

Questions

- What type of lock to get for each operation?
- When should a transaction acquire/release lock?
- We need a ***locking protocol***
 - A set of rules followed by all transactions for requesting and releasing locks

Two phase Locking Protocol (2PL)

- Defines two type of locks:

- *Shared (S) lock*

- *Exclusive (X) lock*

Compatible?	S	X
S	True	False
X	False	False

- Phase 1: Growing Phase

- Each tx must obtain an S (X) lock on an object before reading (writing) it

- Phase 2: Shrinking Phase


- A transaction can not request additional locks once it releases any locks

- Ensures conflict serializability

Implementation

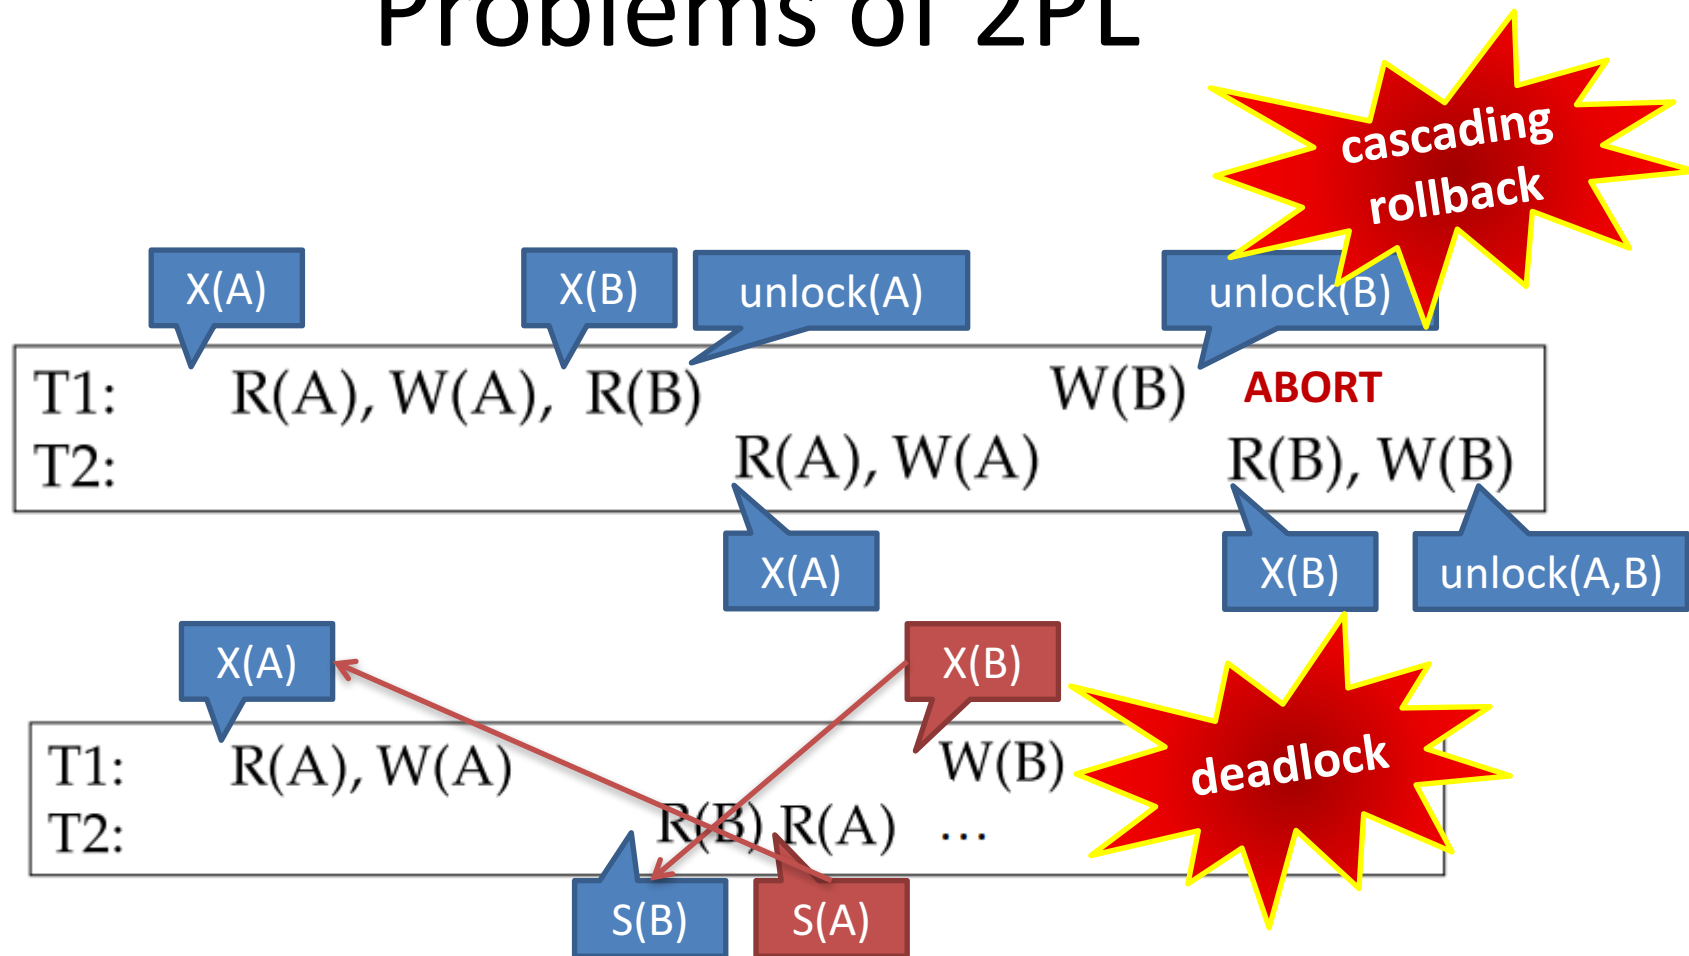
- Lock and unlock requests are handled by the ***lock manager***
 - Shared between concurrency managers
- Lock table entry
 - Number of transactions currently holding a lock
 - Type of lock held
 - Pointer to queue of lock requests
- Locking and unlocking have to be atomic operations

Lock Table



- Implemented as an in-memory hash table indexed on the name of the data item being locked
- New lock request is added to the end of the queue of requests for the data item
- Request is granted if it is compatible with all earlier requests

Problems of 2PL



- **Starvation** is also possible if concurrency control manager is badly implemented

Outline

- Schedules
- Anomalies
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

How to improve 2PL to avoid cascading rollback?

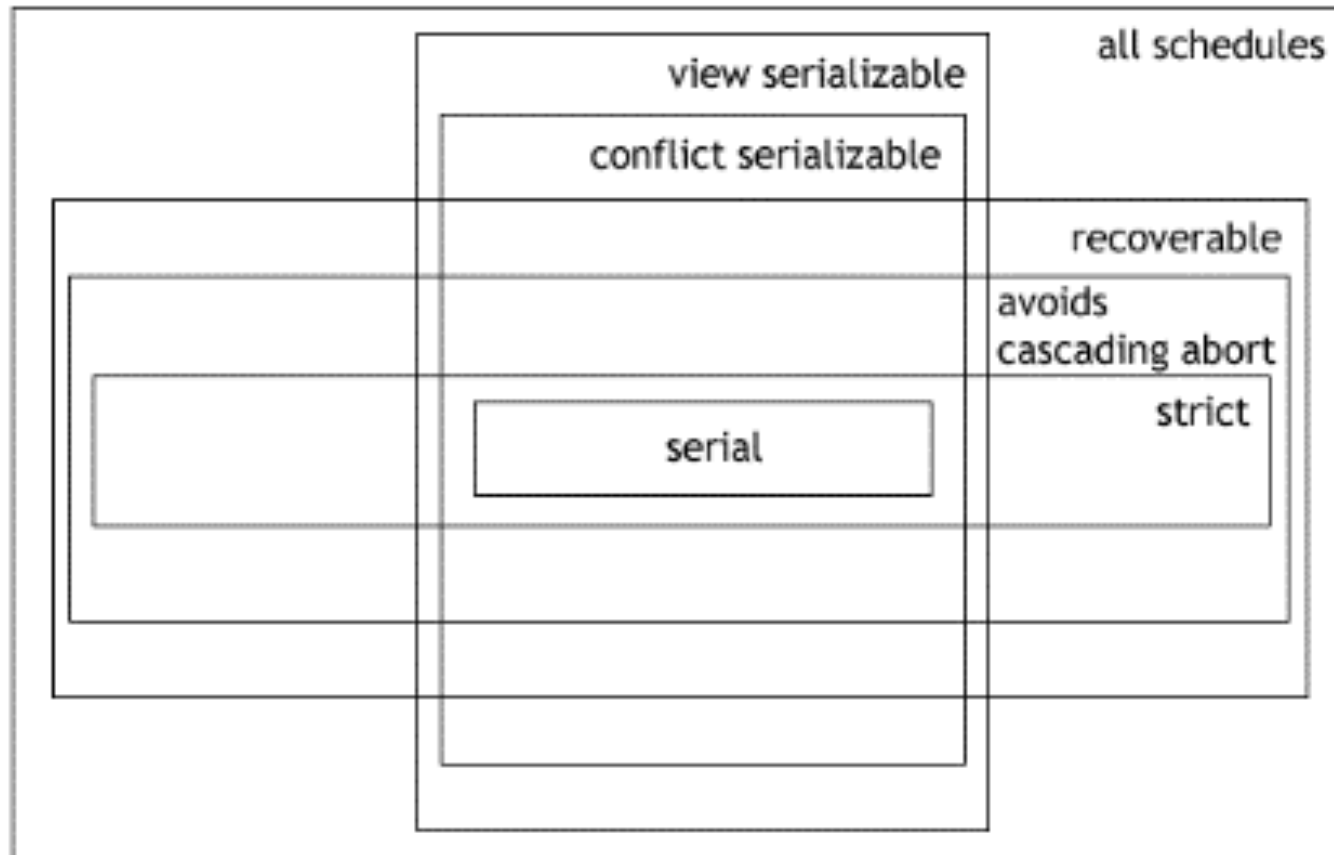
Strict Two-Phase Locking

- S2PL
 1. Each tx obtains locks as in the growing phase in 2PL
 2. But the tx *holds all locks until it completes*
- Allows only serializable and *stric* schedules

Strict Two-Phase Locking

- Definition: A schedule is *strict* iff for any two txs T1 and T2, if a write operation of T1 precedes a conflicting operation of T2 (either read or write), then T1 commits before that conflicting operation of T2
 - Strictness \rightarrow no cascading abort (converse not true)
- Avoids cascading rollback, but still has deadlock

Serializability and Recoverability




Outline


- Schedules
- Anomalies
- Lock-based concurrency control
 - 2 Phase Locking (2PL)
 - Strict 2 Phase Locking (S2PL)
 - Deadlock

Coping with Deadlocks

- Cycle of transactions waiting for locks to be released by each other



- Detection: ***Waits-for*** graph
 - For detecting cycles
- Checked when acquiring locks (or buffers)



Other Techniques (1)

- **Timeout & rollack** (deadlock detection)
 - Assume T_i wants a lock that T_j holds
 1. T_i waits for the lock
 2. If T_i stays on the wait list too long then: T_i is rolled back
- **Wait-die** (deadlock prevention)
 - Assume each T_i has a timestamp (e.g., tx number)
 - If T_i wants a lock that T_j holds
 1. If T_i is older than T_j , it waits for T_j ;
 2. Otherwise T_i aborts

Other Techniques (2)

- **Conservative locking** (deadlock prevention)
 - Every T_i locks **all objects at once** (atomically) in the beginning
 - No interleaving for conflicting txs---performs well only if there is no/very few long txs (e.g., in-memory DBMS)
 - How to know which objects to lock before tx execution?
 - Requires the coder of a stored procedure to specify its read- and write-sets explicitly
 - Does not support ad-hoc queries

References

- Database Design and Implementation, chapter 14. Edward Sciore.
- Database management System 3/e, chapter 16. Ramakrishnan Gehrke.
- Database system concepts 6/e, chapter 15, 16. Silberschatz.
- Derby Developer's Guide: Locking, concurrency, and isolation.
 - <http://db.apache.org/derby/docs/10.9/devguide/cdevconcepts30291.html>
- IBM DB2 document: Locks and concurrency control
 - <http://publib.boulder.ibm.com/infocenter/db2luw/v9r5/index.jsp?topic=%2Fcom.ibm.db2.luw.admin.perf.doc%2Fdoc%2Fc0005266.html>